

TLCTC Framework Glossary — Version 2.5

*Comprehensive definitions and concepts for the Top Level Cyber Threat Clusters framework.
Author: Bernhard Kreinz | Last Updated: 5 Sep 2026*

Implements: TLCTC framework specification **v2.5**. The normative authority for cluster definitions, axioms, and classification rules is the canonical dictionary `json-schemas/layer-1/tlctc-framework.v2.5.json`, reproduced and derived in the v2.5 core paper (`documentation/tlctc-v2.5-core.md`); where an entry here and the canon differ, the canon governs.

This glossary contains all defined terms of the TLCTC framework, organized alphabetically, including the v2.1 boundary extensions, the v2.5 disambiguation rules (R-CHANNEL, R-SUBSTRATE), and industry terminology ("buzzwords") mapped to their correct TLCTC semantic context. Entries carried over from the V2.0/V2.1 whitepaper keep their whitepaper section cross-references. V2.1 additions are marked with (*V2.1*); v2.5 additions with (*v2.5*). Industry terms commonly used in the field are marked with (*Industry Term*). Rule IDs that are no longer part of the v2.5 normative registry are explicitly marked (**Deprecated alias**) — a retired ID keeps its original meaning and is never reused for a different proposition; the sole historical exception (the v2.1 draft R-INTRA numbering) is documented in the R-INTRA entry.

A

Abuse of Functions (#1)

A threat cluster where an attacker misuses the logic, scope, or configuration of existing, legitimate software functions for malicious purposes. This manipulation occurs through standard interfaces using expected input types (data, parameters, configurations, sequence of actions), but in a way that subverts the intended purpose or security controls. Crucially, inputs remain data; no foreign code is introduced or executed. The generic vulnerability is the scope, complexity, or inherent trust placed in legitimate software functions. Classification is governed by the #1 boundary tests in the core paper: if the attacker's success does not require any implementation flaw and instead abuses intended functionality, scope, or configuration via standard interfaces using expected input types, the step **MUST** be classified as #1 Abuse of Functions. (This test was stated as mapping rule R-ABUSE in the v2.0 whitepaper; the ID is a deprecated alias in v2.5.)

Reference: Core paper §4 (#1 boundary tests); formerly R-ABUSE (whitepaper §4.2.5)

Related reading: [AD → Domain Admin → Ransomware cascade](#), [CVE-2026-44578: Next.js WebSocket SSRF](#), [CVE-2020-17103 — patch closed an effect, not a cluster](#), [CrowdStrike 2025 Threat Hunting Report — TLCTC](#), [CrowdStrike 2025 Global Threat Report — TLCTC](#), [The Adoboli Paradox — Cyber vs Operational Risk](#)

Abuse of Rights (v2.5)

An intended action inside an **entitlement** — an access right, a role, a mandate — genuinely conferred by an accountable grantor, but used against its purpose. The envelope is honoured: the system behaved as designed and as authorised, so there is **no generic vulnerability**, **no cluster**,

and no System Risk Event; the consequence chain begins at the Data Risk Event. Register: operational risk, owned by the CRO, the business line and HR, not the CISO. Examples: a badge holder propping open a door, an administrator using genuine root to exfiltrate, a clerk posting a false entry inside their role, an officer approving a payment inside their mandate, a support agent opening a celebrity's record out of curiosity. Controls: segregation of duties, four-eyes, mandate limits, supervision, vetting, purpose auditing.

Why it is not cluster #11. Rights are functions — every authorization check is designed functionality — which is why Abuse of Rights and #1 Abuse of Functions look alike. But rights are the subset of functions that partitions the invocation of all the others, and the taxonomy turns on which side of that partition the action fell. #1 reaches *past* the entitlement using functions the designer left open, and its generic vulnerability is reducible by design (narrow the API, tighten the scope). A granted entitlement is irreducible by design — remove it and you have removed the business. Two exposures, two levers, two owners, two registers.

Boundaries. An insider reaching *outside* their grant is the Attack row and a cluster (the same support agent editing a region parameter to reach a record outside their territory is #1). A **stolen entitlement is never Abuse of Rights**: the grant attaches to the person, not the token, so credential use by a non-grantee is #4 (R-CRED) and the functions reached with it are #1 — an attack wearing an entitlement. Self-enrolment into a population a registration function was not meant for is #1, not Abuse of Rights: nobody with authority conferred anything. The row is a property of the action, not the actor (Axiom IV). The name is **Abuse of Rights**; "Abuse of Access Rights" is deprecated because an access right is only one kind of entitlement.

Reference: Core paper §3.5 (cause-side partition), §6.1 (R-SCOPE); dictionary `cause_side_partition`

Related reading: [Functions and Rights — why one is a threat cluster and the other is not](#), [The Adoboli Paradox — Cyber vs Operational Risk](#)

See also: Cause-Side Partition, Entitlement, Error in Use, R-SCOPE, Abuse of Functions (#1), Operational Risk (OpRisk)

Accessibility (Data Risk Event)

The operational state in which data or resources can be used for their intended purpose by authorized processes (Facility, IT, or Business processes). Loss of Accessibility occurs when data exists but cannot be used—such as encrypted files (ransomware), corrupted data, or permission lockouts. This is distinct from Loss of Availability: ransomware causes Loss of Accessibility (data present but unusable), not Loss of Availability (data gone or unreachable). DRE notation: `Ac` when distinguished from Availability, or `A` as general shorthand covering both.

Actor Archetype

The typical, recurring **attack-sequence pattern** that characterizes how an actor (or an Actor Group) chains TLCTC clusters across incidents — e.g., #9 → #7 → #4 → #1. In the Actor Profile Designer tool the archetype is captured per-actor in the **Archetype** field (the `sequence` attribute in the JSON), and the *Most Common Archetypes* panel aggregates actors that share the same pattern, surfacing dominant attack chains across the dataset. Archetypes are a communication and trend-

watching device — they are never used for cluster classification (Axiom IV). An archetype describes *which clusters an actor tends to prefer* and *in what order it tends to chain them*, not the cluster of any specific attack step. Distinct from [Actor Group](#), which is the generalized actor category (Nation-State, Cybercriminal-Ransomware, ...) and carries a capability vector rather than a sequence pattern.

Reference: §17.3 (Attacker Profiles)

Tool: </tools/actor-profile-designer.html>

Related reading: [CrowdStrike 2024 Threat Hunting Report — TLCTC](#), [CrowdStrike 2025 Threat Report — Strategy & Velocity](#), [Diamond Model × TLCTC — structuring the empty spaces](#), [tlctc-Attacker-Profiling.html](#)

See also: Actor Group, Attacker Profile, Cyber Threat Radar, Axiom IV

Actor Group

A generalized categorization of threat actors used in TLCTC overlays (Attacker Profiles, Tech Enablers Overlay) to group actors by motivation and resourcing level — rather than by named identity. The six reference actor groups published at tlctc.net/tlctc-Attacker-Profiling.html are: **Nation-State**, **Cybercriminal (Ransomware)**, **Cybercriminal (General)**, **Hackivist**, **Insider**, and **Amateur (Script-Kiddie)**. Each actor group carries a baseline **capability vector** across the 10 clusters on the maturity scale (0 = none, 1 = Low, 2 = Medium, 3 = High, 4 = Expert/Champion) — for example, Nation-State scores Expert on #2/#3/#7/#9/#10, whereas Amateur (Script-Kiddie) scores Low across the board. Actor groups are a communication and trend-watching device — they are never used for cluster classification (Axiom IV). An actor group describes *which clusters a class of actors tends to be capable in* and *which emerging tech enablers that class is likely to adopt*, not the cluster of any specific attack step. The typical sequence pattern characteristic of an actor group is captured separately as an [Actor Archetype](#).

Reference: §17.3 (Attacker Profiles), §17.4 (Tech Enablers Overlay)

Tool: </tools/actor-profile-designer.html>

See also: Actor Archetype, Attacker Profile, Tech Enablers Overlay, Cyber Threat Radar, Axiom IV

Attacker Profile (V2.1)

An **informative overlay** on the Cyber Threat Radar that describes a threat actor's (or Actor Group's) observed preferences across the 10 clusters: per-cluster capability scores, preferred cluster sequences, and typical boundary crossings. Profiles may be wrapped in Diamond-Model framing (adversary / capability / infrastructure / victim). Used for hypothesis generation ("likely next steps"), comparative radar views, and targeting analysis against the organizational radar. **Normative guardrails (R-RADAR-6...9):** actors are not clusters, they *use* clusters (Axiom IV); profiles **MUST NOT** redefine cluster meanings or introduce actor-based taxonomy; profiles **SHOULD** be published as probability distributions or pattern frequencies, not deterministic rules; profile scores **SHOULD** be derived from classified Layer 3 attack path instances so provenance traces back to evidence.

Reference: §17.3

Related reading: [CrowdStrike 2024 Threat Hunting Report — TLCTC](#), [CrowdStrike 2025 Threat Report — Strategy & Velocity](#), [Diamond Model × TLCTC — structuring the empty spaces](#), [tlctc-Attacker-Profiling.html](#)

See also: Actor Group, Actor Archetype, Cyber Threat Radar, Axiom IV

Attacker's View

A perspective included in each TLCTC threat cluster definition that describes how the attacker perceives or approaches the exploitation of the specific generic vulnerability. It helps distinguish between clusters by focusing on the attacker's methodology rather than technical implementation details.

Related reading: [LINDDUN vs TLCTC — complementary approaches](#), [TLCTC × Threat Modeling Manifesto](#)

Attack Path

The sequence of applied Attack Vectors in a cyber incident, representing an ordered sequence of Attack Steps describing a complete attack scenario. Basic notation uses #X → #Y → #Z (e.g., #9→#3→#7). Attack paths may include velocity annotations showing the time between steps (e.g., #9→[24h]#4→[12m]#1), domain boundary markers using the || operator, parallel steps, and Data Risk Event tags.

Reference: §4.2.2 (Global Definitions), §11.0 (Path Semantics)

Related reading: [CrowdStrike 2024 Threat Hunting Report — TLCTC](#), [Mandiant M-Trends 2025 — TLCTC](#), [ENISA Threat Landscape 2025 — TLCTC](#), [Same Attack, Four Stories — vendor report comparison](#), [PASTA threat modeling × TLCTC](#), [CKC + ATT&CK + TLCTC — Holy Trinity of Defense](#), [TLCTC × Threat Modeling Manifesto](#), [20 annotated attack paths \(Ransomware, BEC, OT, ...\)](#), [SSDLC for developers — the "S" problem](#), [TLCTC v2.5 monster prompt — CTI & Forensic](#), [Report-to-Radar AI prompt for TLCTC](#), [TLCTC v2.5 monster prompt — SOC & Detection](#), [TLCTC v2.1 AI analysis prompt — teach any LLM](#), [TLCTC v2.5 monster prompts — index](#)

Attack Path Notation

The standardized format for describing cyber attack sequences using TLCTC clusters. Format uses: → for sequential steps, + for parallel execution, [time] for temporal intervals, and |[context] |[Source→@Target]|| for domain boundaries. Example: #9→[24h]#4→[12m]#1 |[dev]|[@Vendor→@Org]|| →[weeks]#10.2→[0s]#7.

Related reading: [CVE-2020-17103 — patch closed an effect, not a cluster](#), [20 annotated attack paths \(Ransomware, BEC, OT, ...\)](#), [Dual-layer notation — TLCTC-XX.YY enumeration](#), [TLCTC v2.5 monster prompt — CTI & Forensic](#)

Attack Sequence Schema (V2.0)

The JSON schema that defines the required structure for documenting attack path instances. The schema ensures all documented attacks follow a consistent format including metadata (sequence_id, attack_title, framework_version), temporal transitions (delta_t_value, delta_t_unit, velocity_class), responsibility spheres, and cluster mappings. File format: tlctc-attack-sequence-schema.json.

Attack Step (V2.0)

A single attacker action or event that exploits exactly **one generic vulnerability** in a specific context. Each Attack Step MUST map to exactly one TLCTC cluster (per Axiom VI).

Reference: §4.2.2 (Global Definitions), §4.2.8 (Minimal Classification Procedure)

Attack Vector

The specific path or method used by an attacker to gain unauthorized access to a target system. In the TLCTC framework, each distinct attack vector is a distinct initiating method defined by the **initial generic vulnerability targeted** (per Axiom VII). The vector label MUST be based on cause, not outcome.

Reference: §4.2.2 (Global Definitions), Axiom VII (§2)

Related reading: [Evolving VERIS — replace Action axis with TLCTC](#)

Attack Velocity (Δt) (V2.0)

The temporal dimension of cyber risk representing the **time interval** between two adjacent Attack Steps in an attack path. For an edge $\#X \rightarrow \#Y$, the value $\Delta t(X \rightarrow Y)$ represents the elapsed time between step $\#X$ and step $\#Y$ in the described scenario. Δt is an edge property attached to the sequence operator, not to steps. Attack velocity matters defensively because it bounds the time a detection-and-response control has to act between adjacent steps (see Detection Coverage Score). Categorized into four velocity classes: Latent/Slow (days to months), Medium (hours), Fast (minutes), and Realtime (seconds/milliseconds).

Reference: §12.0 (Definitions), §12.1 (Measurement Model), §12.2 (Notation)

Related reading: [CrowdStrike 2025 Threat Hunting Report — TLCTC](#), [CrowdStrike 2025 Threat Report — Strategy & Velocity](#), [GTIG AI Threat Tracker \(May 2026\) — TLCTC](#), [MITRE ATT&CK & STIX × TLCTC V2.0 — implementation guide](#), [IEC 62443 × TLCTC v2.0 — industrial cybersecurity](#), [FAIR × TLCTC — enhanced quantitative risk](#), [ISO/SAE 21434 × TLCTC V2.0 — automotive](#), [ISO 27000 × TLCTC — name vs game](#), [OCTAVE × TLCTC v2.0 — causal taxonomy](#), [Enhancing CVE records with TLCTC v2.1](#), [EU Cybersecurity Act \(CSA\) × TLCTC V2.0](#), [EU cyber regulation needs a common taxonomy](#), [DORA TLPT × TLCTC V2.1 — boundary & velocity](#), [Why ORX must rethink the "cyber event"](#), [TLCTC classification decision tree V2.0/V2.1](#), [Agentic AI as consequence amplifier \(right side of Bow-Tie\)](#), [Quantum & AI — new magic, same 10 threats](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#), [tlctc-attack-velocity.html](#), [TLCTC v2.5 monster prompt — SOC & Detection](#)

Availability (Data Risk Event)

The technical state in which data or resources exist and can be reached by the infrastructure. Availability answers the question: "Does the data exist and can the system access it?" A resource is available if it is present, can be enumerated, and is technically accessible to the system—regardless of whether it can be meaningfully used. For example, encrypted files remain available (they exist on disk) even when rendered unusable by ransomware. Loss of Availability occurs when data is deleted, storage fails, or systems go offline. DRE notation: Av when distinguished from Accessibility, or A as general shorthand covering both.

Axiom

A foundational premise that defines what terms mean and what kinds of statements are allowed in TLCTC. Axioms are non-negotiable constraints on interpretation that force methodological consequence and prevent logical shortcuts and category errors. These foundational principles must be accepted to validate and effectively use the TLCTC framework. TLCTC defines ten axioms organized into four groups: Scope (I–II), Separation (III–V), Classification (VI–VIII), and Sequence (IX–X).

Reference: §2 (Axioms and Assumptions)

Related reading: [End of Semantic Diffusion — Kuhn & TLCTC](#), [Logical foundations of TLCTC](#), [Why exactly ten? — TLCTC architecture](#), [10×10×10 — fun fact about TLCTC](#), [Why TLCTC does not need the "Hazard"](#), [The Control Fixation Reflex](#), [Understanding cyber threats — a common language](#)

A (continued — Industry Terms)

Adware (*Industry Term*)

A type of malware that delivers unwanted advertisements, often bundled with legitimate software. In TLCTC: maps to #7 Malware — the software environment's designed execution capability is abused to run foreign advertising code. Adware that arrives via a compromised software package may involve a #10 → #7 sequence.

See also: Malware (#7), Supply Chain Attack (#10)

AI / AGI / ASI (Positioning in TLCTC)

Artificial Intelligence, Artificial General Intelligence, and Artificial Super Intelligence occupy three distinct roles in the TLCTC framework:

- **As an IT system:** AI is exposed to the same 10 threat clusters as any other IT system (software + hardware).
- **As a tool:** AI enhances the capabilities of both threat actors AND defenders.
- **As AGI/ASI:** Would become a powerful threat actor OR defender in its own right.

AI systems do not create new threat clusters — they are subject to the existing 10 clusters and may amplify their exploitation.

Reference: V1.9.1 Clarifications

Related reading: [GTIG AI Threat Tracker \(May 2026\) — TLCTC](#), [MITRE ATLAS × TLCTC integration](#), [MITRE ATT&CK for ML \(AML\) × TLCTC](#), [NIST AI RMF + MITRE ATLAS × TLCTC](#), [Agentic AI as consequence amplifier \(right side of Bow-Tie\)](#), [Agentic AI under the microscope — TLCTC](#), [Quantum & AI — new magic, same 10 threats](#), [AI conversation deep dive on TLCTC](#), [Grok AI 4.1 — independent TLCTC validation](#), [G7 SBOM-for-AI — control fixation critique](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#), [Emerging technologies as threat enablers](#), [TLCTC v2.1 AI analysis prompt — teach any LLM](#)

Amplification Attack (*Industry Term*)

A flooding technique where an attacker sends small requests to third-party services (e.g., NTP, DNS, memcached) that respond with disproportionately large replies directed at the victim. In TLCTC: maps to #6 Flooding Attack — the primary mechanism is volume exceeding finite capacity. The abuse of the amplification service itself may additionally involve #1 Abuse of Functions.

See also: Flooding Attack (#6), DDoS

ARP Spoofing (*Industry Term*)

A technique where an attacker sends falsified ARP (Address Resolution Protocol) messages on a local network, linking the attacker's MAC address to a legitimate IP address. In TLCTC: the ARP spoofing itself is #1 Abuse of Functions (abusing a legitimate network protocol without exploiting an implementation flaw). ARP spoofing typically leads to a Man in the Middle position, making the full sequence #1 → #5.

Reference: V1.9.1 Buzz-Word Refinement (#1)

See also: Man in the Middle (#5), DNS Spoofing, SSL Stripping

B

Bounded Δt

A minimum or maximum bound for Δt derived from known constraints when precise timestamps are unavailable. Notation: $\Delta t < 15m$ (upper bound), $\Delta t > 15m$ (lower bound), $\Delta t = 10m..20m$ (range).

Reference: §12.0.3, §12.2.3

Related reading: tlctc-attack-velocity.html

Bow-Tie Model

A risk model that represents risk as a structure with five elements: Threats (left side), Preventive Controls (left side), Central Event (knot), Mitigating Controls (right side), and Consequences (right side). TLCTC is anchored in the Bow-Tie model to enforce strict separation between cause and effect in cyber risk analysis. The model enforces temporal causality, prevents confusion between threats and outcomes, enables precise control placement, and reveals attack sequences as causal chains. The central event "Loss of Control" serves as the pivot point between threat realization and potential consequences.

Reference: §6 (The Bow-Tie Anchor), §6.1 (Structure and Vocabulary)

Related reading: [Basel operational risk × TLCTC Bow-Tie](#), [Why ORX must rethink the "cyber event"](#), [TLCTC × Threat Modeling Manifesto](#), [Generic vulnerabilities — software & hardware failure](#), tlctc-bow-tie-causality.html, blog-tlctc-two-layer-framework.html, [Why TLCTC does not need the "Hazard"](#), [Agentic AI as consequence amplifier \(right side of Bow-Tie\)](#)

Bridge Cluster

A TLCTC cluster whose generic vulnerability **inherently** enables crossing into (or leveraging over) a different domain's control regime. Bridge clusters are: #8 Physical Attack, #9 Social Engineering, and #10 Supply Chain Attack.

Reference: §4.2.2 (Global Definitions), §5.1.4, §5.2 (Topology Classification)

Related reading: [TLCTC classification decision tree V2.0/V2.1](#), [Topology of cyber attacks — Bridge vs Internal](#)

Bridge Step

A step-level instance of a bridge cluster that crosses a specific domain boundary. When a bridge step crosses responsibility spheres, the boundary SHOULD be recorded in path notation via the domain boundary operator `||[context][@Source→@Target]||`.

Reference: §4.2.2 (Global Definitions), §5.1.6

BGP Hijacking (*Industry Term*)

A technique where an attacker manipulates Border Gateway Protocol routing tables to redirect internet traffic through attacker-controlled infrastructure. In TLCTC: the BGP manipulation itself is #1 Abuse of Functions (abusing legitimate protocol functionality, not an implementation flaw). BGP hijacking typically leads to a Man in the Middle position, making the full sequence #1 → #5.

Reference: V1.9.1 Buzz-Word Refinement (#1)

See also: ARP Spoofing, DNS Spoofing, Man in the Middle (#5)

Botnet (*Industry Term*)

A network of compromised devices ("bots" or "zombies") controlled by an attacker, typically used to amplify attack capabilities. In TLCTC: the botnet itself is the result of prior #7 Malware infections on each device. When a botnet is used for flooding, the attack maps to #6 Flooding Attack. The full lifecycle may involve: #9 → #7 (social engineering delivering malware to build the botnet), then #6 (coordinated flooding from the botnet).

See also: Flooding Attack (#6), Malware (#7), DDoS

Brute-Force Attack (*Industry Term*)

A method of systematically trying all possible credential combinations (passwords, PINs, encryption keys) to gain unauthorized access. In TLCTC: maps to #4 Identity Theft — the attacker is attempting to derive and use credentials to impersonate a legitimate identity. The generic vulnerability is the insufficient binding, at the point of authentication, between a presented credential and the authentic holder: the endpoint cannot distinguish a guessing attacker from the legitimate holder when controls such as account lockout, rate-limiting, or strong-password requirements are absent. (Weak credential *storage/protection* is a separate, acquisition-side concern that classifies to the enabling cluster, not #4.)

Reference: V1.9.1 Buzz-Word Refinement (#4)

See also: Identity Theft (#4), Password Spraying, Credential / Identity Artifact

Buffer Overflow (*Industry Term*)

A class of implementation flaw where a program writes data beyond the boundaries of allocated memory, potentially allowing an attacker to execute arbitrary code or crash the application. In TLCTC: maps to #2 Exploiting Server or #3 Exploiting Client depending on whether the vulnerable component is in a server role (accepting inbound requests) or client role (consuming external responses) per R-ROLE. Buffer overflows create an unintended data→code transition.

Reference: V1.9.1 §Definitions (#2, #3), Buzz-Word Refinement

See also: Exploiting Server (#2), Exploiting Client (#3), Implementation Flaw

Business Risk Event (BRE)

A discrete, observable business-level event on the consequence side of the Bow-Tie model, triggered by a Data Risk Event or by a preceding BRE. Examples include regulatory notification obligations, service outage declarations, media coverage, customer churn, and regulatory fines. BREs may **chain**: each BRE can trigger subsequent BREs, forming a variable-length consequence sequence (SRE → DRE → BRE₁ → BRE₂ → ... → BRE_n). Each BRE→BRE transition has its own Δt representing a detection and intervention window. An organization's Risk Appetite determines at which point a BRE is designated as the terminal **Business Impact (BI)** — BI is a role a BRE can hold, not a separate event category.

Reference: §6.3.1 (The Consequence Chain), V1.9.1 §The Anatomy of Risk

Related reading: [Evolving VERIS — replace Action axis with TLCTC](#), [blog-cyber-bow-tie-business-risk-event-chain.html](#), [Agentic AI as consequence amplifier \(right side of Bow-Tie\)](#), [Propagated Controls — Rule of Propagation](#), [TLCTC v2.5 monster prompt — Regulators & Standards](#), [TLCTC+ for NCSCs & CERTs — national reporting](#)

See also: System Risk Event (SRE), Data Risk Event (DRE), Business Impact (BI), Event Chain, Consequences

Business Impact (BI)

A **role** assigned to the terminal Business Risk Event in a consequence chain — the BRE beyond which further causal decomposition is no longer operationally useful for a given organization. BI is not a separate event category; it is the point in the BRE chain where an organization's **Risk Appetite** boundary is reached. What constitutes BI for one organization may be a mid-chain BRE for another: a €50K regulatory fine may be terminal impact for a startup but a mid-chain event for a multinational. The BI designation is therefore context-dependent and organization-specific.

Reference: §6.3.1 (The Consequence Chain)

See also: Business Risk Event (BRE), Risk Appetite, Event Chain

BxIs (Base Level Indicators)

The lowest level of indicators that still make operational sense, representing metrics at the operational level directly translated into measurable values. Part of the hierarchical KxI framework (KRIs, KCIs, KPIs).

C

BEC (Business Email Compromise) / CEO Fraud (*Industry Term*)

A social engineering attack where an adversary impersonates a senior executive or trusted business partner (often via compromised or spoofed email) to trick employees into transferring funds, revealing sensitive information, or taking other harmful actions. In TLCTC: the manipulation of the human target maps to #9 Social Engineering. If the attacker uses compromised email credentials, the email access maps to #4 Identity Theft. Typical sequences: #9 → #1 (social engineering leading to function abuse, e.g., wire transfer) or #4 → #9 → #1 (stolen credentials enabling impersonation for social engineering leading to function abuse).

Reference: V1.9.1 Buzz-Word Refinement (#9 — "CEO Fraud", "Invoice Manipulation Fraud (BEC Fraud)")

Related reading: [20 annotated attack paths \(Ransomware, BEC, OT, ...\)](#)

See also: Social Engineering (#9), Whaling, Phishing

Call-Level Mapping Rule

A TLCTC classification principle for function-call-level threat analysis:

- **Parameter tampering**, unauthorized function selection, or misuse of valid functions without executing foreign code → always #1 Abuse of Functions.
- **Presentation of identity artifacts** at call time (e.g., stolen API keys, session tokens, cookies, Kerberos tickets) to impersonate a subject → always #4 Identity Theft.

This rule prevents overlap between logic misuse (#1) and identity presentation/use (#4) at the function call level, where the caller acts as "client" and the called function as "server".

Reference: V1.9.1 §Concept Applicability (At Function Call Level)

See also: Abuse of Functions (#1), Identity Theft (#4)

CAPEC (Common Attack Pattern Enumeration and Classification) (*Industry Term*)

A MITRE-maintained dictionary of known attack patterns, each describing a method of exploiting known weaknesses. In TLCTC: CAPEC patterns are operational-level detail that map to the strategic-level TLCTC clusters. CAPEC complements CWE (weaknesses) and CVE (specific vulnerabilities) in the conceptual hierarchy, and TLCTC proposes bridging from NIST to the extended MITRE world (ATT&CK, CWE, CAPEC, CVE) through the 10 clusters.

Reference: V1.9.1 §Standardizing Strategic Cybersecurity

See also: CWE, CVE, MITRE ATT&CK, Techniques (TTPs)

Capacity Exhaustion

Degradation or denial of service caused **primarily** by volume or intensity exceeding finite resources. Resources include: bandwidth, CPU cycles, memory, storage, database connections, API quotas, thread/process pools, file handles. Maps to #6 Flooding Attack.

Reference: §4.2.2 (Global Definitions), R-FLOOD (§4.2.5)

Cause-Side Partition (v2.5)

The partition of the cause side of any risk event into four rows by three questions asked in strict order: **Is there an actor? Did they intend it? Were they entitled?**

Actor	Intent	Entitlement	Row	Register	Clusters
No	—	—	Failure / external event	OpRisk	none
Yes	No	—	Error in Use	OpRisk	none
Yes	Yes	Yes	Abuse of Rights	OpRisk	none — no SRE; chain starts at the DRE
Yes	Yes	No	Attack	Cyber	the ten TLCTC clusters; SRE recorded

Entitlement is asked only once intent is present (an unentitled actor who did not intend the outcome is Error in Use, not an attack), and *entitled* means entitled, not permitted. The partition is a property of the action, not of the actor (Axiom IV): an outsider with no grant and an insider outside their grant land in the same Attack row. Only the Attack row is in TLCTC scope; the ten clusters classify its steps and only its steps (R-SCOPE). Third party is a modifier over all four rows, carried by the domain-boundary operator, not a fifth row. Decision procedure: (1) actor? no → failure; (2) intended? no → Error in Use; (3) entitlement covering *this* action? yes → Abuse of Rights; (4) implementation flaw required? yes → #2/#3 per R-ROLE, no → #1.

Reference: Core paper §3.5; dictionary `cause_side_partition`

See also: Abuse of Rights, Error in Use, Entitlement, R-SCOPE, System Risk Event (SRE), Operational Risk (OpRisk)

Central Event

In the TLCTC Bow-Tie model: the **System Risk Event (SRE)**, which since v2.5 has two types at one altitude — **System Compromise / Loss of Control** (the point at which an actor achieves unauthorized control over the system's behavior, privileges, data, or trust relationships, sufficient to pursue attack objectives; the only type the ten clusters reach) and **System Failure / Loss of Function** (no actor holds anything). The central event is positioned before outcomes.

Reference: §6.3 (Central Event)

Client-Role Component

A component that **consumes external responses, content, or state** relative to the attacker. The component is in "client role" for the specific interaction being classified.

Reference: §4.2.2 (Global Definitions), R-ROLE (§4.2.5)

Client-Server Relationship

A fundamental principle (Axiom II) stating that every software system interaction — networked or intra-system — is based on client-server or caller-called function interaction at various levels. A network is not a precondition: a syscall, hypercall, or IPC call establishes the same caller-called relation as a remote protocol exchange. The relationship is contextual: the entity requesting a service is the "client," and the entity providing that service is the "server". Roles can be dynamic and change depending on interaction context, particularly across protection ring boundaries.

Coder

A development role focused on implementation and craftsmanship, responsible for writing functional, efficient code according to established patterns, implementing specific security controls at the code level, and following secure coding practices. Primary responsibility for addressing threat clusters #2, #3, and implementation details of #4, #5, and #7. Contrasts with the Programmer role which focuses on architecture and strategy.

Related reading: [Taming SonarQube with TLCTC \(Secure-Dev pt 2\)](#), [Programmer vs Coder in TLCTC \(Secure-Dev pt 1\)](#), [Adding the Developer's View to TLCTC](#)

Consequences

In the Bow-Tie model: what results after the central event, including technical and business impact (event chains). Consequences are on the right (effect) side of the Bow-Tie and are recorded as Data Risk Events. Consequences are NOT threat categories.

Reference: §6.1 (Bow-Tie Structure), §6.4 (What TLCTC Does NOT Classify)

Control

A security measure implemented to mitigate threats, reduce vulnerabilities, or minimize the impact of security incidents. In the TLCTC framework, controls are organized using NIST CSF functions (Identify, Protect, Detect, Respond, Recover) and mapped to specific threat clusters. Controls are categorized as Local Controls (protecting specific systems) or Umbrella Controls (protecting groups of systems).

Control Design Effectiveness

An evaluation of whether a control, as conceived and structured, is theoretically capable of achieving its objective if it operates as intended. Assesses the control's capability to address the identified risk within its specific threat cluster.

Related reading: [D3FEND × TLCTC — the missing axis](#), [TLCTC Control Matrix Manager](#), [Control Matrix for SME & Private \(TLCTC v2.3\)](#)

Control Failure

A deviation from a control objective or lack of effectiveness. Control failure is control-risk and MUST NOT be treated as a threat category (Axiom V). Risk structure remains: Threat → Event/Incident →

Consequences; controls influence likelihood and impact but do not define the threat cluster. Distinguished from the actual risk event itself (Axiom IV).

Reference: Axiom V (§2), §6.2 (Rule 3)

Related reading: [CVE-2020-17103 — patch closed an effect, not a cluster](#), [Logical impossibility of control-first regulation](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#), [The Control Fixation Reflex](#), [The Audit Trap — compliance ≠ security](#)

Control Objective

The specific aim or purpose that a control is intended to achieve, defining what the control should accomplish in terms of risk mitigation for a particular threat cluster. Each control aligns with a single, clear objective.

Related reading: [NIST SP 800-218 \(SSDF\) × TLCTC](#), [SABSA × TLCTC — architecture × threat ontology](#), [D3FEND × TLCTC — the missing axis](#), [Logical impossibility of control-first regulation](#), [G7 SBOM-for-AI — control fixation critique](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#), [The Control Fixation Reflex](#), [TLCTC Control Matrix Manager](#), [Propagated Controls — Rule of Propagation](#), [Control Matrix for SME & Private \(TLCTC v2.3\)](#), [Capability-based planning via 10×\(6×2\) matrix](#), [The Audit Trap — compliance ≠ security](#)

Control Operational Effectiveness

An evaluation of whether a control is actually working as designed in practice, examining if the control is being executed correctly and consistently over time to meet its objective. May vary depending on the nature of the threat cluster (e.g., controls for Malware #7 may never achieve 100% due to detection latencies).

Related reading: [TLCTC Control Matrix Manager](#)

Credential / Identity Artifact

Any secret, token, key, or session artifact that enables authentication or authorization decisions. Examples include: passwords, PINs, passphrases, API keys, bearer tokens, OAuth/OIDC tokens, SAML assertions, session cookies, session identifiers, private keys, client certificate keys, Kerberos tickets, SSH keys, hardware token seeds/OTPs, biometric templates (when used as authenticators).

Reference: §4.2.2 (Global Definitions), Axiom X (§2)

Credential Acquisition

The act of obtaining, capturing, exposing, deriving, or forging a credential/identity artifact. Credential acquisition maps to the **enabling cluster**—the generic vulnerability that made the acquisition possible.

Reference: §4.2.2 (Global Definitions), R-CRED (§4.2.5), Axiom X (§2)

Credential Application

The act of presenting, using, replaying, or leveraging a credential to authenticate and operate as an identity **other than the presenter's own**. So scoped, credential application **MUST** always map to

#4 Identity Theft. Authenticating as an identity the target system itself issued to the presenter through a designed enrolment function is authentication as self, not credential application in this sense — see the self-issued proviso under R-CRED.

Reference: §4.2.2 (Global Definitions), R-CRED (§4.2.5), Axiom X (§2)

Credential Forgery

The act of creating a credential without possessing the legitimate secret. If forgery succeeds due to an implementation flaw (e.g., weak signing algorithm, missing validation, predictable tokens), the forgery step maps to #2 or #3 per R-ROLE. The subsequent use of the forged credential maps to #4.

Reference: §4.2.2 (Global Definitions), R-CRED (§4.2.5)

CVE (Common Vulnerabilities and Exposures)

A standardized identifier for publicly known cybersecurity vulnerabilities. In the TLCTC framework, CVEs are mapped to generic vulnerabilities and their corresponding threat clusters to enable consistent threat classification and control implementation.

Related reading: [CVE-2020-17103 — patch closed an effect, not a cluster](#), [Enhancing CVE records with TLCTC v2.1](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#), [Harbor Registry × TLCTC Cyber Threat Radar](#)

Cyber Bow-Tie

The specific application of the Bow-Tie Model to cyber risk management, with the 10 Top Level Cyber Threat Clusters on the cause side, "Loss of Control" or "System Compromise" as the central event, and Data Risk Events and Business Risk Events on the consequence side. Enables structured cyber risk register development and event chain analysis.

Cyber Incident

An actual security breach or system compromise that has occurred, representing the materialization of a cyber risk event where control over IT systems or persons has been lost due to one or more of the 10 Top Level Cyber Threat Clusters.

Cyber Risk

The probability of occurrence of a cyber event in which control over IT systems or persons is lost due to one or more of the 10 Top Level Cyber Threat Clusters, leading (via event chains) to consequential damage (impact). Cyber risks are a subset of operational risks (OpRisk).

Related reading: [ENISA Gap Analysis — TLCTC](#), [NIST NICE × TLCTC — workforce capabilities](#), [FAIR × TLCTC — enhanced quantitative risk](#), [ISO 27000 × TLCTC — name vs game](#), [OCTAVE × TLCTC v2.0 — causal taxonomy](#), [EU regulation \(NIS2/DORA/CRA\) vs TLCTC](#), [Why DORA will fail regarding cyber risks](#), [TLCTC vs 30+ standards & regulations](#), [Basel operational risk × TLCTC Bow-Tie](#), [Cyber Resilience Act \(CRA\) — TLCTC pain points & fixes](#), [Generic vulnerabilities — software & hardware failure](#), [TLCTC — the missing link between strategy and ops](#), [The Audit Trap — compliance ≠ security](#), [The Adoboli Paradox — Cyber vs Operational Risk](#), [TLCTC v2.5 monster prompt — CISO](#)

[& Risk, The Risk Appetite Conversation Your Board Isn't Having](#), [TLCTC — the missing link \(brief\)](#), [Strategic risk management implementation guide v2.1](#)

Cyber Risk Event

A potential occurrence that could lead to a system breach or compromise. Distinguished from Cyber Incidents (which have already occurred) and Data Risk Events (which are consequences). The central event in the Cyber Bow-Tie model is the **System Risk Event (SRE)** — "Loss of Control" or "System Compromise".

Related reading: [GDPR vs NIS2 — different trigger points](#), [Basel operational risk × TLCTC Bow-Tie](#), [Why ORX must rethink the "cyber event"](#)

See also: System Risk Event (SRE)

CWE (Common Weakness Enumeration) (*Industry Term*)

A community-developed list of common software and hardware weakness types maintained by MITRE. CWE categorizes the underlying flaws, bugs, or errors (weaknesses) that enable vulnerabilities to exist — e.g., CWE-89 for SQL Injection weakness, CWE-119 for buffer overflow weakness. In the TLCTC conceptual hierarchy: **Weakness (CWE)** → **Specific Vulnerability (CVE)** → **Generic Vulnerability (TLCTC)** → **Threat Cluster (#1–#10)**. CWE provides granular weakness taxonomy at the code level for developers; TLCTC operates at the strategic level by grouping all resulting vulnerabilities into 10 generic vulnerability categories.

Related reading: [CVE-2026-21510: Windows Shell SmartScreen bypass](#), [CVE-2020-17103 — patch closed an effect, not a cluster](#), [TLCTC × CWE — interactive explorer](#), [TLCTC × CWE — context is king](#), [TLCTC × CWE — advanced explorer](#), [MITRE CWE needs a taxonomic reboot](#), [CWE-514 \(covert channel\) → TLCTC #8](#), [SonarQube × TLCTC — secure dev integration \(pt 3\)](#), [Taming SonarQube with TLCTC \(Secure-Dev pt 2\)](#), [SSDLC for developers — the "S" problem](#), [Adding the Developer's View to TLCTC](#), [TLCTC v2.5 monster prompt — DevSecOps](#)

See also: Vulnerability, Weakness, CVE, Generic Vulnerability

Cyber Threat Radar

A standard visualization methodology for communicating threat posture, change over time, and comparative exposure across the 10 TLCTC clusters. The radar is a **communication surface**, **not a classification device** — every position on it is driven by evidence classified under the Section 4 grammar.

Structure:

- **Spokes (10):** one per cluster (#1–#10), fixed order; the invariant layout is what makes radars from different organizations, sectors, or snapshots directly comparable.
- **Zones (4, outer → inner):** Latent → Low → Medium → High. Radial distance encodes current exposure or impact.
- **Sectors (configurable):** angular subdivisions within each spoke representing responsibility scopes. Default alignment is to Layer 2 responsibility spheres (@Org, @Customers, @Vendor /

@External). At state/national scope, sectors may encode critical-infrastructure sectors (energy, finance, healthcare, etc.).

- **Bubbles:** individual threat instances, sized by significance.
- **Movement indicators:** ▲ rising / ▼ falling exposure vs. the previous snapshot.

Scales: organizational view (enterprise-internal), sector / line-of-business view, and state-level / national view. The spoke layout is invariant across all three — only the sector axis changes, which is what enables cross-view comparison.

Normative rules (R-RADAR-1...5): spoke assignment MUST use the Section 4 grammar; multi-cluster attack paths MUST be rendered as multiple bubbles or a single bubble plus a separate Layer 3 path; zone placement SHOULD follow a disclosed scoring method; snapshots SHOULD be dated and use the same method for movement indicators to be meaningful; aggregated sector / national radars MUST use identical spoke definitions and SHOULD disclose the combination rule (max / average / weighted).

Reference: §17.1–17.2

Related reading: [CrowdStrike 2024 Threat Hunting Report — TLCTC](#), [ENISA Threat Landscape 2025 — TLCTC](#), [GTIG AI Threat Tracker \(May 2026\) — TLCTC](#), [Emerging technologies as threat enablers](#), [Harbor Registry × TLCTC Cyber Threat Radar](#), [Report-to-Radar AI prompt for TLCTC](#)

See also: Attacker Profile, Tech Enablers Overlay, Responsibility Sphere, Layer 2

Tool: </tools/radar-tlctc-app.html>

Command Injection (*Industry Term*)

An attack where an attacker injects operating system commands into an application that passes user input to a system shell. In TLCTC: maps to #2 Exploiting Server or #3 Exploiting Client per R-ROLE — this is an implementation flaw (failure to sanitize input) that creates an unintended data→code transition. The execution of the injected commands constitutes Foreign Executable Content, so the full sequence is typically #2 → #7 or #3 → #7.

See also: Exploiting Server (#2), SQL Injection, Foreign Executable Content (FEC)

Cross-Site Scripting (XSS) (*Industry Term*)

A class of implementation flaw where an application includes untrusted data in web output without proper validation or encoding, allowing attacker-controlled scripts to execute in a victim's browser. In TLCTC:

- **Stored/Reflected XSS** (server fails to encode output): #2 Exploiting Server — the flaw is in server-side code.
- **DOM-Based XSS** (client script processes data unsafely): #3 Exploiting Client — the flaw is in client-side code.

In both cases, the XSS creates an unintended data→code transition via an implementation flaw.

Reference: V1.9.1 §Definitions (#2, #3), Buzz-Word Refinement

See also: Exploiting Server (#2), Exploiting Client (#3), Implementation Flaw

D

Data Processing Pathways

The four distinct paths that data can follow during an attack, each mapping to specific TLCTC clusters:

1. Data → Data (#1 only) — Pure data manipulation without code execution
2. Data → Function Invocation → Foreign Code Execution (#1 → #7) — Function abuse enabling code execution
3. Data → Exploit Code via Implementation Flaw (#2 or #3) — Unintended data→code transition
4. Data → Foreign Code via Designed Execution Capability (#7 only) — Intended execution capability

Data Risk Event (DRE)

An outcome event describing **Loss of Confidentiality (C)** (data stolen / unauthorized access), **Loss of Integrity (I)** (data modified / unauthorized changes), or **Loss of Availability/Accessibility (A)** (data gone or unreachable, or data present but unusable). Data Risk Events **MUST** be recorded separately from cluster steps, **MUST NOT** be used as threat categories, and **MUST NOT** change the cluster classification of the step that preceded them. Notation: [DRE: C], [DRE: I], [DRE: A], or combinations. The type codes form a **refinement tree** (core §7.6; dictionary `data_risk_events`), not a flat list: **I** **MAY** be refined into **Ii** (Incorrect State — correspondence or completeness fails; the content is wrong) or **If** (Misattributed State — provenance or attribution fails; the record claims an origin it does not have, and its content may be perfectly accurate); **A** **MAY** be refined into **Av** (Unavailable State — data gone or unreachable) or **Ac** (Inaccessible State — data present but unusable). A parent code stays legal whenever the refinement is unknown or irrelevant. **Stopping rule:** a property is split only where the resulting states are distinguishable by inspecting the record itself, never by the story of how it got there — which admits exactly these two splits, refuses every split of **C**, and refuses any split by cause (Axiom III). Refinements are read off the record, never off the actor; there is no "manipulated" code. Examples: ransomware encryption = [DRE: Ac]; data deletion = [DRE: Av]; a value keyed wrongly on purpose = [DRE: Ii]; an entry posted under someone else's session = [DRE: If]; a transcript showing a command that never ran, written by code standing in for the harness = [DRE: Ii, If]; distinction unknown = [DRE: I] / [DRE: A].

Reference: §4.2.2 (Global Definitions), §6.2 (Rule 2), §11.5.3

Related reading: [Chaos Ransomware — TLCTC forensic](#), [Evolving VERIS — replace Action axis with TLCTC](#), [LINDDUN vs TLCTC — complementary approaches](#), [Enhancing CVE records with TLCTC v2.1](#), [GDPR vs NIS2 — different trigger points](#), [TLCTC classification decision tree V2.0/V2.1](#), [TLCTC+ for NCSCs & CERTs — national reporting](#)

Data vs Code Boundary

A normative classification principle: Domain-specific expressions (e.g., SQL, LDAP, XPath, GraphQL, template syntax, configuration languages) are treated as **data** unless they directly cause **FEC**

execution via a general-purpose execution engine.

Reference: §4.2.2 (Global Definitions)

Delta t (Δt) (V2.0)

Symbol representing the time interval between threat cluster transitions in an attack sequence. See Attack Velocity.

Designed Execution Capability

The environment's **intended** capability to load, interpret, or execute program content. This is the generic vulnerability exploited by #7 Malware. Examples: OS loaders, script interpreters, macro engines, browser JS engines, module loaders, container/virtualization runtimes.

Reference: §4.2.2 (Global Definitions), §4.1 (#7 Definition)

Related reading: [The File Type Fallacy — extension blocklists](#), [GovCERT-CH blocked filetypes × TLCTC](#)

Detection Coverage Score (DCS) (V2.0)

A strategic Key Performance Indicator (KPI) measuring **detection timing adequacy** relative to Attack Velocity. Formula: $DCS = (\text{Mean Time to Detect}) / (\text{Attack Velocity } \Delta t)$. It answers one question: is detection latency shorter than the attacker's progression window between two adjacent steps?

- **Score < 1.0:** Detection is faster than the adversary's progression (the response window exists)
- **Score > 1.0:** Adversary completes the step before detection (no response window)

Example: If a ransomware group moves from #4 to #1 in 10 minutes and your SIEM alerts in 15 minutes, $DCS = 15/10 = 1.5$, indicating systematic blindness requiring automation rather than analyst intervention.

Scope note. Despite the historical name, DCS measures timing adequacy, not coverage: it assumes the relevant activity is detectable at all. A detector with 10-second MTTD but low detection probability does not have good coverage merely because Δt is 60 seconds; detection probability and rule coverage must be assessed separately (see the application paper, Part B).

Related reading: [The Commit Is the CVE — silent fixes & the patch-gap collapse](#)

Developer's View

A perspective included in each TLCTC threat cluster definition that provides guidance on secure development practices specific to preventing that cluster. Encompasses both Programmer (architectural) and Coder (implementation) responsibilities.

Related reading: [Programmer vs Coder in TLCTC \(Secure-Dev pt 1\)](#), [SSDLC phase-by-phase reference — TLCTC v2.1](#), [Adding the Developer's View to TLCTC](#), [TLCTC v2.5 monster prompt — DevSecOps](#)

Domain

A set of assets governed by a coherent control regime (policies, monitoring, enforcement, and accountability). Domains may be technical, organizational, or socio-technical. Examples: cyber/IT domain, physical security domain, human decision domain, vendor development domain, cloud provider control-plane domain.

Reference: §4.2.2 (Global Definitions), §5.1.1

Domain Boundary

A point where responsibility spheres or control regimes change. Crossing a domain boundary means the attack moves from one set of applicable controls to a different set.

Reference: §4.2.2 (Global Definitions), §5.1.3

Related reading: [Topology of cyber attacks — Bridge vs Internal](#)

Domain Boundary Operator (||) (V2.0)

Notation: `||[context][@Source→@Target]||`. Used to explicitly mark where an attack path crosses responsibility spheres. The operator SHOULD accompany bridge cluster steps (#8, #9, #10) and MAY be used with any step that crosses a domain boundary. The context describes the transition type (e.g., [dev], [idp], [update]) and the arrow shows the direction of trust crossing. The boundary test: "If removing the third-party trust link would stop the step from succeeding, #10 belongs there". Enables precise mapping of responsibility shifts and supply chain attack analysis.

Reference: §4.2.2 (Global Definitions), §11.3 (Domain Boundary Operator), §5.3

Related reading: [MITRE ATT&CK & STIX × TLCTC V2.0 — implementation guide](#), [IEC 62443 × TLCTC v2.0 — industrial cybersecurity](#), [FAIR × TLCTC — enhanced quantitative risk](#), [ISO/SAE 21434 × TLCTC V2.0 — automotive](#), [Enhancing CVE records with TLCTC v2.1](#), [EU cyber regulation needs a common taxonomy](#), [DORA TLPT × TLCTC V2.1 — boundary & velocity](#), [TLCTC classification decision tree V2.0/V2.1](#), [Topology of cyber attacks — Bridge vs Internal](#), [TLCTC+ for NCSCs & CERTs — national reporting](#)

DAST (Dynamic Application Security Testing) (Industry Term)

A testing methodology that analyzes a running application by simulating attacks against it to identify security vulnerabilities. In TLCTC: DAST is a **preventive control** (IDENTIFY/PROTECT) primarily targeting #2 Exploiting Server and #3 Exploiting Client by discovering implementation flaws in deployed applications.

See also: SAST, Control, Exploiting Server (#2), Exploiting Client (#3)

DDoS (Distributed Denial of Service) (Industry Term)

An attack where multiple compromised systems (typically a botnet) simultaneously flood a target with traffic or requests. In TLCTC: maps to #6 Flooding Attack — the primary mechanism is volume or intensity exceeding finite capacity. The "distributed" aspect describes the attack infrastructure, not the generic vulnerability. Sub-types include SYN Flood, UDP Flood, HTTP Flood, ICMP Flooding,

NTP/DNS Amplification, and Slowloris. The underlying botnet is the result of prior #7 Malware infections.

Reference: V1.9.1 Buzz-Word Refinement (#6)

See also: Flooding Attack (#6), Botnet, SYN Flood, Amplification Attack

Defense-in-Depth (*Industry Term*)

A security strategy employing multiple layers of controls so that if one layer fails, another provides protection. In TLCTC: defense-in-depth means implementing controls at multiple points along potential attack paths — both Local Controls (specific systems) and Umbrella Controls (groups of systems) — and across NIST CSF functions (Identify, Protect, Detect, Respond, Recover) for each relevant threat cluster. The interplay between different threat clusters (e.g., #9 potentially circumventing #4 controls) necessitates a holistic, defense-in-depth approach.

Reference: V1.9.1 §The Anatomy of Risk

See also: Local Controls, Umbrella Controls, Control

Directory Traversal (*Industry Term*)

An attack where an attacker manipulates file path references (e.g., using `../` sequences) to access files or directories outside the intended scope. In TLCTC: maps to #2 Exploiting Server — an implementation flaw in how the server-side code handles file path input, enabling unauthorized access. The flaw is in server-side source code (failure to validate/sanitize path references).

Reference: V1.9.1 Buzz-Word Refinement (#2)

See also: Exploiting Server (#2), Implementation Flaw

DNS Spoofing (*Industry Term*)

A technique where an attacker corrupts DNS resolution to redirect traffic to attacker-controlled infrastructure. In TLCTC: the DNS spoofing itself is #1 Abuse of Functions (abusing legitimate DNS protocol functionality without exploiting a code flaw). DNS spoofing typically leads to a Man in the Middle position, making the full sequence #1 → #5. Note: if the attacker exploits an implementation flaw in a DNS server (e.g., cache poisoning via a code bug), the initial step maps to #2 Exploiting Server instead.

Reference: V1.9.1 Buzz-Word Refinement (#1)

See also: ARP Spoofing, BGP Hijacking, Man in the Middle (#5)

Domain Squatting (*Industry Term*)

Registering domain names similar to legitimate ones (typosquatting, homograph attacks) to deceive users into visiting attacker-controlled websites. In TLCTC: the domain registration is infrastructure setup (not itself a cluster step). When used to harvest credentials via fake login pages, the attack maps to #9 Social Engineering (luring the user) → #4 Identity Theft (using the stolen credentials). When used to deliver malware, the sequence may be #9 → #3 or #9 → #7.

Reference: V1.9.1 Buzz-Word Refinement (#4)

See also: Phishing, Social Engineering (#9), Typosquatting

DORA (Digital Operational Resilience Act) (*Industry Term*)

An EU regulation establishing ICT risk management requirements for financial entities. Like NIS2, DORA emphasizes incident reporting but lacks a unified threat categorization system. The TLCTC framework addresses this gap by providing a standardized taxonomy that supports compliance with DORA's requirements for threat classification and incident reporting. DORA requirements can be mapped as Regulatory Trigger Points in the TLCTC event chain model.

Related reading: [CVE-2020-17103 — patch closed an effect, not a cluster](#), [EU cyber regulation needs a common taxonomy](#), [EU regulation \(NIS2/DORA/CRA\) vs TLCTC](#), [Why DORA will fail regarding cyber risks](#), [DORA TLPT × TLCTC V2.1 — boundary & velocity](#), [TLCTC vs 30+ standards & regulations](#), [Logical impossibility of control-first regulation](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#), [TLCTC v2.5 monster prompt — Regulators & Standards](#)

See also: NIS2, Regulatory Trigger Point, E_n Event Notation

Drive-By Download (*Industry Term*)

An attack where malware is automatically downloaded and potentially executed when a user visits a compromised or malicious website, typically by exploiting a browser or plugin vulnerability. In TLCTC: maps to #3 Exploiting Client (browser vulnerability exploited) → #7 Malware (payload execution). If the user was lured to the site via social engineering, the full sequence is #9 → #3 → #7.

Reference: V1.9.1 §F (Edge-Case Resolution)

See also: Exploiting Client (#3), Malware (#7), Watering Hole Attack

Dual-Use Tool

A legitimate administrative utility that can be used for both legitimate administrative purposes and malicious activities when invoked by an attacker. Examples include PowerShell, PsExec, WMI, and remote administration tools. In TLCTC: invocation/abuse of the tool may be #1 (if no implementation flaw is exploited), while the actual execution of attacker-controlled FEC through that tool is #7, resulting in a #1 → #7 sequence.

Reference: §4.2.5 (R-EXEC, LOLBAS Clarification)

E

Edge (in attack path)

A transition between two adjacent Attack Steps, represented by the sequence operator →. Δt (Attack Velocity) is an edge property.

Reference: §11.1 (Sequence Operator), §12.0.2

Entitlement (v2.5)

What an accountable grantor actually conferred on a person for an action: an access right, a role, or a mandate, held genuinely and bounded by its purpose and scope. **Entitled is not permitted.** Permission is what the access-control system happens to return; entitlement is what a grantor conferred. The two are supposed to coincide and frequently do not, and that gap is where the classification of a technical act can depend on a governance document — the framework refuses to pretend a question about authorization can be answered without consulting the authority. **An entitlement attaches to the person, never to the token:** an attacker holding a credential was never a grantee, so no action they take is inside any envelope (#4 per R-CRED, then #1). The entitlement test is the authorization-side counterpart of R-CRED's authentication-side test (is the system deceived about who is authenticating?): both ask about the truth of a relation, never about the identity of the claimant.

Reference: Core paper §3.5; dictionary `cause_side_partition.entitlement`

See also: Abuse of Rights, Cause-Side Partition, R-SCOPE, R-CRED

Error in Use (v2.5)

The second row of the cause-side partition: an actor — any actor — producing an outcome they did not intend. Entitlement is not asked: a visitor who trips a breaker is Error in Use, not #8; an administrator who mistypes a command and drops a table is Error in Use, not #1. Operational risk, no cluster, no attacker's view. Error in Use can enter any altitude directly and can produce any Data Risk Event type — a misdirected email is `c`, an accidental purge is `Av`, a value keyed wrongly is `Ii` — and, where it breaks a system, it produces a System Failure, not a System Compromise.

Reference: Core paper §3.5; dictionary `cause_side_partition`

See also: Cause-Side Partition, Abuse of Rights, System Failure, Operational Risk (OpRisk)

Estimated Δt

An approximate Δt value derived from partial evidence when precise timestamps are unavailable. Notation: $\Delta t \sim 15m$.

Reference: §12.0.3, §12.2.3

E_n Event Notation (Regulatory) (V2.0)

A numbered event sequence notation used to map attack chains to regulatory compliance triggers:

- **E1:** System Compromise / Loss of Control (the central Bow-Tie event)
- **E2:** Data Risk Event (e.g., PII exposure — GDPR trigger)
- **E3a, E3b, ...:** Compliance violation events (e.g., GDPR breach notification, NIS2 incident report)

The subscript (a, b, etc.) distinguishes parallel regulatory branches triggered by the same upstream event. Different regulations trigger at different points: GDPR Art. 33 triggers at E2 (Data Risk Event involving PII), while NIS2 Art. 23 triggers at E1 (Significant Incident). See also: Event Chain Length, RS Container.

Event Chain Length (V2.0)

The number of causal events between the initial incident (E1) and a regulatory trigger point (E3x). Shorter chains mean compliance clocks start sooner. Example: NIS2 path (E1→E3b) = 2 events with 24h early warning requirement; GDPR path (E1→E2→E3a) = 3 events with 72h notification timeline. Understanding chain length helps CISOs structure incident response playbooks to meet multiple regulatory timelines from a single incident.

Exploit Code

Foreign code that targets specific vulnerabilities to modify software behavior, creating an UNINTENDED data→code transition. Used in #2 Exploiting Server and #3 Exploiting Client. Distinguished from Malware Code which operates within expected execution paths.

Exploiting Client (#3)

A threat cluster where an attacker targets and leverages implementation flaws within any component acting in a client role (requesting/processing data from a server or resource). These vulnerabilities allow manipulation of client behavior or unauthorized access using Exploit Code, often when the client interacts with malicious content. The generic vulnerability is that client-side implementation flaws enable unintended behavior. "Implementation" is substrate-neutral — the flawed logic may live in application source code, firmware, microcode, or hardware logic; per R-SUBSTRATE, the location of a flaw never determines the cluster, the role of the flawed component does.

Related reading: [CVE-2025-21333 revisited: #2 after all \(Kernel's Role pt 2\)](#), [Apache 2.4.67 — 11 CVEs decomposed](#), [CVE-2026-21510: Windows Shell SmartScreen bypass](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#)

Exploiting Server (#2)

A threat cluster where an attacker targets and leverages implementation flaws within a component acting in a server role. These vulnerabilities allow manipulation of server behavior or unauthorized access using Exploit Code, forcing a data→code transition where exploit code executes as new, foreign code in the server context. The generic vulnerability is that server-side implementation flaws enable unintended behavior. "Implementation" is substrate-neutral — the flawed logic may live in application source code, firmware, microcode, or hardware logic; per R-SUBSTRATE, the location of a flaw never determines the cluster, the role of the flawed component does.

Related reading: [Calif M5: #2 → #2 \(Kernel's Role pt 1\)](#), [Apache 2.4.67 — 11 CVEs decomposed](#), [CVE-2026-31431 \(Copy Fail\): Linux kernel AF_ALG](#), [CVE-2026-35414: 15-year-old OpenSSH cert flaw](#), [CVE-2026-46300 \(Fragnesia\): Linux kernel XFRM](#), [Verizon DBIR 2025 — TLCTC](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#)

EDR (Endpoint Detection and Response) (Industry Term)

A category of security tools that monitor endpoint devices for suspicious activity and provide automated response capabilities. In the TLCTC velocity model, EDR is a critical control for the **Fast Velocity Class** (minutes) where automated containment is necessary because human analyst response times are insufficient. EDR is particularly relevant for controlling #3 Exploiting Client and #7 Malware at the endpoint level.

See also: Fast Velocity Class, SIEM, SOAR

Event Chain

A causal sequence where one outcome event triggers subsequent events, following the consequence chain **SRE** → **DRE** → **BRE***. The chain cascades from the System Risk Event (central event) through Data Risk Events to one or more Business Risk Events. BREs may themselves chain ($BRE_1 \rightarrow BRE_2 \rightarrow \dots \rightarrow BRE_n$), with each transition having its own Δt representing a detection and intervention window where all six NIST CSF functions apply. Example: System Compromise (SRE) → Data Breach involving PII (DRE [C]) → GDPR notification obligation (BRE₁) + NIS2 incident report (BRE₂) → Regulatory fine (BRE₃). Understanding event chains is critical for designing Respond/Recover controls and regulatory compliance workflows.

Reference: §6.3.1 (The Consequence Chain), V1.9.1 §Data Risk Event Types, §Clarification on Central Event Position

Related reading: [blog-cyber-bow-tie-business-risk-event-chain.html](#), [Propagated Controls — Rule of Propagation](#), [Strategic risk management implementation guide v2.1](#)

See also: System Risk Event (SRE), Data Risk Event (DRE), Business Risk Event (BRE), Business Impact (BI), E_n Event Notation, RS Container, Propagated PR

Evil Maid Attack (*Industry Term*)

A physical attack where an adversary with brief unsupervised access to a device (e.g., left in a hotel room) tampers with it — installing hardware keyloggers, modifying boot loaders, or extracting encryption keys. In TLCTC: the physical access step maps to #8 Physical Attack. Subsequent technical steps map to their respective clusters: installing a keylogger = #8 → #7, extracting credentials from the device = #8 → #4.

Reference: V1.9.1 Buzz-Word Refinement (#8)

See also: Physical Attack (#8), USB Baiting

F

Fast Velocity Class (V2.0)

A velocity classification where attack progression occurs within minutes. Typical threat clusters: #3 (Exploiting Client), #2 (Exploiting Server). Control strategy requires automated containment and EDR blocking, as human analyst response times are insufficient. Example: Drive-by downloads, browser exploits, RCE exploitation.

Flooding Attack (#6)

A threat cluster where an attacker intentionally overwhelms system resources or exceeds capacity limits through a high volume of requests, data, or operations, leading to disruption, degradation, or denial of service for legitimate users. The generic vulnerability is the finite capacity limitations

inherent in any system component (network bandwidth, CPU, memory, storage, database limits, application quotas, API rate limits, process/thread pools). Outcome is typically Loss of Availability.

Related reading: [Apache 2.4.67 — 11 CVEs decomposed](#)

Fileless Execution / Fileless Malware (*Industry Term*)

An attack technique where malicious code executes entirely in memory without writing traditional files to disk, often using legitimate system tools (PowerShell, WMI, .NET reflection) as execution vehicles. In TLCTC: fileless execution still maps to #7 Malware — the FEC definition explicitly includes in-memory execution, interpreted code, and reflective loading with no "on-disk" requirement. The invocation of the legitimate tool to enable fileless execution may be #1 Abuse of Functions, making the typical sequence #1 → #7.

See also: Foreign Executable Content (FEC), Living Off the Land / LOLBAS, Dual-Use Tool

Foreign Executable Content (FEC)

Attacker-controlled (or otherwise untrusted) program text or bytes that are **interpreted, loaded, or executed** by a **general-purpose execution engine** in the target environment. Includes attacker-controlled commands fed into interpreters. FEC execution includes in-memory (fileless) execution, interpreted code, macro execution, and reflective loading—no "on-disk" requirement exists.

Reference: §4.2.2 (Global Definitions)

Related reading: [The File Type Fallacy — extension blocklists](#), [GovCERT-CH blocked filetypes × TLCTC](#)

Framework Layer (V2.0)

The static, universal component of the TLCTC JSON architecture containing threat cluster definitions, generic vulnerabilities, data risk events, bow-tie model principles, attack path notation rules, and framework axioms. Defined in `tlctc-framework.json`. Changes rarely (only during framework evolution) and serves as the common language that all organizations reference. Contrasts with the Intelligence Layer which contains dynamic, incident-specific data.

G

Generic Vulnerability

The single root-level vulnerability category defining a cluster — the **strategic-level attack surface** towards a specific class of threats. For every generic vulnerability, there is exactly one TLCTC cluster (per Axiom VI). The generic vulnerability is what the threat cluster targets: it is the exposed surface that enables the attack vector and through which the attack path proceeds. Generic vulnerabilities are stable across technologies and implementations, persisting regardless of specific IT system types, software implementations, or evolving attack techniques. All specific vulnerabilities (CVEs) are instances of a generic vulnerability; all generic vulnerabilities map to exactly one threat cluster.

The 10 generic vulnerabilities (attack surfaces) are: functional scope/trust (#1), server-side implementation flaws (#2), client-side implementation flaws (#3), identity-artifact binding (#4), lack of end-to-end communication protection (#5), finite capacity limitations (#6), designed execution capability (#7), physical accessibility (#8), human psychological factors (#9), and third-party trust dependencies (#10).

Reference: §4.2.2 (Global Definitions), §4.2.8 (Step 2), Axiom VI (§2)

Related reading: [Mandiant M-Trends 2025 — TLCTC](#), [ENISA Threat Landscape 2025 — TLCTC](#), [ENISA Gap Analysis — TLCTC](#), [Same Attack, Four Stories — vendor report comparison](#), [NIST threat definitions — structural gap](#), [22 NIST definitions of "threat" — TLCTC](#), [Evolving VERIS — replace Action axis with TLCTC](#), [LINDDUN vs TLCTC — complementary approaches](#), [PASTA threat modeling × TLCTC](#), [IEC 62443 × TLCTC v2.0 — industrial cybersecurity](#), [Diamond Model × TLCTC — structuring the empty spaces](#), [ISO/SAE 21434 × TLCTC V2.0 — automotive](#), [ISO 27000 × TLCTC — name vs game](#), [OCTAVE × TLCTC v2.0 — causal taxonomy](#), [MITRE CWE needs a taxonomic reboot](#), [EU Cybersecurity Act \(CSA\) × TLCTC V2.0](#), [Cyber Resilience Act \(CRA\) — TLCTC pain points & fixes](#), [TLCTC × Threat Modeling Manifesto](#), [End of Semantic Diffusion — Kuhn & TLCTC](#), [Logical foundations of TLCTC](#), [Why exactly ten? — TLCTC architecture](#), [Generic vulnerabilities — software & hardware failure](#), [blog-tlctc-two-layer-framework.html](#), [Why TLCTC does not need the "Hazard"](#), [Agentic AI under the microscope — TLCTC](#), [AI conversation deep dive on TLCTC](#), [TLCTC — the missing link between strategy and ops](#), [Tactics evolve, 10 threats are constant](#), [Tactics evolve, 10 threats are constant — TLCTC](#), [TLCTC for everyone — the Blind Spot method](#), [Programmer vs Coder in TLCTC \(Secure-Dev pt 1\)](#), [SSDLC phase-by-phase reference — TLCTC v2.1](#), [TLCTC v2.5 monster prompt — DevSecOps](#), [Understanding cyber threats — a common language](#), [TLCTC — the missing link \(brief\)](#)

See also: Vulnerability, Attack Vector, Attack Surface Analysis

GOVERN (GV)

The governance function in NIST CSF 2.0, operating at a strategic level to establish the overall cybersecurity risk management framework. GOVERN controls are "assurance controls" that create structure and context for other functions, including setting risk appetite, defining roles and responsibilities, and establishing policies. While GOVERN oversees threat categorization in the risk register, GV controls themselves don't directly counter specific threats but provide the strategic foundation for comprehensive risk management.

H

HTTP Flood (*Industry Term*)

An application-layer denial of service attack that overwhelms a web server with seemingly legitimate HTTP requests. In TLCTC: maps to #6 Flooding Attack — the primary mechanism is volume exceeding finite capacity at the application layer. Distinguished from implementation-flaw-based DoS (which maps to #2 or #3 per R-FLOOD).

See also: Flooding Attack (#6), DDoS, Slowloris, SYN Flood

I

ICMP Flooding (*Industry Term*)

A network-layer denial of service attack that overwhelms a target with ICMP echo request (ping) packets. In TLCTC: maps to #6 Flooding Attack — volume exceeding finite network capacity.

See also: Flooding Attack (#6), DDoS

Identity Theft (#4)

A threat cluster where an attacker presents or uses authentication credentials (passwords, tokens, keys, session identifiers, biometrics) they do not authentically control to impersonate a legitimate identity (human or technical). The generic vulnerability is insufficient binding, at the point of authentication, between a presented credential and the authentic holder of the identity it claims. Credential storage and transmission are prevention controls that reduce acquisition; acquisition-side failures classify to the enabling cluster (#2/#5/#7/#8), not to #4.

Critical distinction: Credentials have dual operational nature:

- **Acquisition/Exposure:** When credentials are obtained through another cluster (e.g., #2 SQL injection, #5 MitM, #7 keylogger, #9 Phishing), map to the enabling cluster (Loss of Confidentiality consequence)
- **Use/Application:** The subsequent *use* of acquired credentials—regardless of acquisition method—always maps to #4 Identity Theft (Loss of Control / system compromise event)

Non-Overlap Rule: Credential acquisition maps to the enabling threat cluster; credential use always maps to #4.

Related reading: [AD → Domain Admin → Ransomware cascade](#), [CVE-2026-35414: 15-year-old OpenSSH cert flaw](#), [CrowdStrike 2025 Threat Hunting Report — TLCTC](#), [CrowdStrike 2025 Global Threat Report — TLCTC](#), [CrowdStrike 2025 Threat Report — Strategy & Velocity](#), [Verizon DBIR 2025 — TLCTC](#), [blog-MFABypass.html](#)

Implementation Defect (Availability Context)

A flaw in code logic, parsing, memory handling, or resource handling that causes crash, hang, or degradation when triggered—**without** requiring volume/intensity to exceed normal capacity. Includes algorithmic complexity weaknesses (e.g., ReDoS). Maps to #2 or #3 per R-ROLE, not #6.

Reference: §4.2.2 (Global Definitions), R-FLOOD (§4.2.5)

Implementation Flaw

A defect in implemented logic (logic, parsing, memory handling, resource handling) enabling unintended behavior when triggered — whether that logic is realized in application source code, firmware, microcode, or hardware description logic (substrate-neutral per R-SUBSTRATE). Implementation flaws are exploited by #2 Exploiting Server (server-role) or #3 Exploiting Client (client-role).

Reference: §4.2.2 (Global Definitions), §4.1 (#2 and #3 Definitions)

Insider Threat (*Industry Term*)

Not a TLCTC category. The industry label mixes two rows of the cause-side partition that answer to different control regimes: an insider acting **inside** their entitlement against its purpose is **Abuse of Rights** (operational risk, no cluster, no SRE), while an insider reaching **outside** their entitlement is the **Attack** row and classifies to a cluster exactly as an outsider would (Axiom IV: the row is a property of the action, not the actor). A programme built around "insider threat" therefore procures endpoint and network telemetry for events that never pass through a compromise, and segregation-of-duties controls for events that do. TLCTC splits the label at the entitlement envelope (R-SCOPE).

See also: Abuse of Rights, Cause-Side Partition, Entitlement, R-SCOPE

Intelligence Layer (V2.0)

The dynamic component of the TLCTC JSON architecture containing specific attack instances, software versions & CVEs, timeline & actor TTPs, domain boundaries, and impact assessments. Changes constantly as new incidents occur. Each incident is documented in its own JSON file following the attack sequence schema format: `[incident-id]-attack-path.json`. Enables worldwide threat intelligence sharing while maintaining consistency through reference to the static Framework Layer.

Internal Cluster

A TLCTC cluster that operates primarily **within the software domain's** attack surfaces, without inherently crossing to a different responsibility sphere. Internal clusters are: #1 through #7.

Reference: §4.2.2 (Global Definitions), §5.1.5, §5.2 (Topology Classification)

Insecure Deserialization (*Industry Term*)

A class of implementation flaw where an application deserializes untrusted data without proper validation, potentially allowing arbitrary code execution or object manipulation. In TLCTC: maps to #2 Exploiting Server or #3 Exploiting Client depending on the role of the vulnerable component per R-ROLE. The deserialization flaw creates an unintended data→code transition via an implementation defect.

Reference: V1.9.1 Buzz-Word Refinement (#2, #3)

See also: Exploiting Server (#2), Exploiting Client (#3), Implementation Flaw

Intra-System Boundary Operator (|...|) (V2.1)

Notation: `[[type][@from→@to]]`. Used to annotate boundary crossings **within a single host or system**, such as sandbox escapes, privilege escalations, process boundary violations, and VM escapes. Uses single pipe delimiters to distinguish from the inter-sphere Domain Boundary Operator (`(|...|)`). Defined boundary types: `sandbox`, `privilege`, `process`, `hypervisor`. The `memory` type is reserved and **MUST NOT** be used (R-INTRA-9). Intra-system boundaries are observability annotations and never change cluster classification (R-INTRA-7). Example: #3 `[[sandbox][@renderer→@os]]` — browser exploit escaping renderer sandbox.

Reference: §11.3.6 (Intra-System Boundary Operator)

J

JSON Architecture (V2.0)

The standardized data structure for threat intelligence sharing in TLCTC (introduced in v2.3), consisting of four complementary JSON files:

1. **tlctc-framework.json**: Core framework definitions (universal, rarely updated)
2. **tlctc-responsibility-spheres.json**: Domain boundary definitions (customizable, occasionally updated)
3. **tlctc-attack-sequence-schema.json**: Validation schema for attack instances (universal, rarely updated)
4. **[incident]-attack-path.json**: Specific attack instances (per-incident, constantly updated)

This architecture separates universal framework definitions from specific attack instances, enabling machine-readable, validated, consistent threat intelligence exchange worldwide.

K

KCI (Key Control Indicator)

A metric that measures the operational performance of security controls, verifying that intended actions are taken at the appropriate frequency. KCIs provide insights on the ability to apply correct controls correctly, highlighting process weaknesses and tool effectiveness. Example: "Frequency of patch deployments per day" or "Scan verification of implemented patches" for a control requiring critical systems to be patched within 24 hours.

Related reading: [CVE-2020-17103 — patch closed an effect, not a cluster](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#)

KPI (Key Performance Indicator)

A measurable value demonstrating the outcome and performance of security processes in reaching security objectives. KPIs must be time-based and reflect effectiveness over time. Example: "Average time to restore critical services to full operation within a 4-hour window." In TLCTC, the Detection Coverage Score (DCS) serves as a strategic detection-timing KPI (see its entry for scope caveats).

KRI (Key Risk Indicator)

A leading indicator demonstrating the potential for a future cyber threat. KRIs show possible risks before a threat occurs and must be observed in a meaningful timeframe. Example: "Number of unpatched critical vulnerabilities older than 7 days" indicates how processes handle critical vulnerabilities, helping identify, understand, and prioritize security efforts to prevent incidents.

Keylogger (*Industry Term*)

Malware (software) or a hardware device that records keystrokes to capture sensitive information such as passwords, credit card numbers, or other data. In TLCTC: a software keylogger maps to #7

Malware (foreign code executing via the environment's designed execution capability). A hardware keylogger maps to #8 Physical Attack (physical device installation). Captured credentials that are subsequently used map to #4 Identity Theft. Typical sequence: #7 → #4 (software keylogger) or #8 → #4 (hardware keylogger).

Reference: V1.9.1 Buzz-Word Refinement (#7, #8)

See also: Malware (#7), Physical Attack (#8), Identity Theft (#4)

Kill Chain (*Industry Term*)

A model describing the stages of a cyber attack from reconnaissance through exploitation to objective completion (originally Lockheed Martin Cyber Kill Chain). In TLCTC: the kill chain concept is complementary to — but not a classification criterion for — the TLCTC framework. TLCTC classifies each step by the generic vulnerability exploited (cause-oriented), while kill chain models describe the phase of the attack (process-oriented). Attack paths in TLCTC notation (#X → #Y → #Z) naturally encode the kill chain progression without requiring a separate phase model.

Reference: V1.9.1 §F (Oversimplification)

Related reading: [The Kill Chain Fallacy — process is not taxonomy](#), [CKC + ATT&CK + TLCTC — Holy Trinity of Defense](#)

See also: Attack Path, Sequence, MITRE ATT&CK

KxI Framework

The integrated hierarchical framework of Key Risk Indicators (KRIs), Key Control Indicators (KCIs), and Key Performance Indicators (KPIs), providing a practical mechanism to operationalize the 10 Top Level Cyber Threat Clusters. Each threat cluster has associated KRI, KCI, and KPI values for managing cyber risk and measuring overall cybersecurity program performance. Base Level Indicators (BxIs) represent the lowest operational level that still makes sense.

L

Latent/Slow Velocity Class (V2.0)

A velocity classification where attack progression occurs over days to months. Typical threat clusters: #10 (Supply Chain), #7 (APT Implants). Control strategy focuses on log retention and threat hunting, as detection windows are extended. Example: Supply chain compromises with long dwell times, persistent APT campaigns prioritizing stealth over speed.

Related reading: tlctc-attack-velocity.html

Living Off the Land / LOLBAS (Living Off the Land Binaries and Scripts)

An attack technique using only software functions and binaries already present on a (potentially compromised) system, invoked with legitimate inputs/parameters, without introducing foreign code initially. Legitimate system binaries are used to execute attacker-controlled content. In TLCTC: the invocation of the legitimate binary may be #1 (if no implementation flaw is exploited), while the

execution of attacker-controlled content through it is #7. The sequence #1 → #7 applies. Examples: Using cmd.exe, PowerShell, WMI, or Task Scheduler to execute attacker-controlled scripts.

Reference: §4.2.5 (R-EXEC, LOLBAS Clarification)

Lateral Movement (*Industry Term*)

The techniques an attacker uses to progressively move through a network after initial compromise, seeking higher-value targets and expanded access. In TLCTC: lateral movement is not a single cluster — it is an attack path composed of multiple sequential steps. Typical lateral movement sequences include: #4 → #1 (using stolen credentials to access another system's functions), #4 → #7 (deploying malware on additional systems using stolen credentials), or #1 → #7 (abusing legitimate remote tools to execute code on other systems). Each step in lateral movement maps to its own cluster based on the generic vulnerability exploited.

Reference: V1.9.1 §Bridging Strategy and Operations, §E (Real World Examples)

Related reading: [AD → Domain Admin → Ransomware cascade](#)

See also: Attack Path, Sequence, Identity Theft (#4)

Local Controls

Security measures implemented directly on or for specific IT systems. Distinguished from Umbrella Controls which protect groups of systems. Local controls are essential for systems that cannot be fully protected by umbrella controls (e.g., exposed systems, "Patient Zero" entry points).

Loss of Availability (LoA)

A Data Risk Event outcome where data or resources are gone or unreachable — the resource no longer exists or cannot be technically accessed by the infrastructure. From the attacker's perspective: "Data gone / system down". Examples: deletion, storage failure, system offline, network unreachable. DRE notation: **Av** (refined) or **A** (general). Not to be confused with Loss of Accessibility (data present but unusable, e.g., ransomware encryption → **Ac**).

Loss of Accessibility (LoAc)

A Data Risk Event outcome where data or resources exist and can be reached by the infrastructure, but cannot be used for their intended purpose by authorized processes. From the attacker's perspective: "Data locked / unusable". Examples: ransomware encryption, data corruption, permission lockout. DRE notation: **Ac** (refined) or **A** (general). Not to be confused with Loss of Availability (data gone or unreachable → **Av**).

Loss of Confidentiality (LoC)

A Data Risk Event outcome where an attacker gains unauthorized access to data. From the attacker's perspective: "Data stolen". This describes what bad thing happens, not how it happens. Various threat clusters can lead to this outcome depending on the mechanism used (e.g., #2 via SQL injection, #5 via MitM eavesdropping).

Disambiguation: In TLCTC, **LoC** always means Loss of Confidentiality — a *consequence-side* Data Risk Event. "Loss of Control" is a distinct concept: the Bow-Tie

central event, always abbreviated **SRE** (System Risk Event), never "LoC". See **System Risk Event (SRE)**.

Loss of Control / System Compromise

The **Compromise type** of the System Risk Event (SRE) — the central event in the Cyber Bow-Tie model — representing the point at which an actor achieves unauthorized control over a system's behavior, privileges, data, or trust relationships. Since v2.5 the SRE has two types at one altitude: Compromise (this entry; the only type the ten clusters reach) and System Failure (loss of function, no actor). This serves as the pivot point between threat realization (cause) and potential consequences (effect). The SRE is the first event in the consequence chain: **SRE** → **DRE** → **BRE***. Some attacks may have delayed data risk events (creating a detection window), while others lead to immediate data risk events. Examples: A server exploit (#2) enabling remote code execution leading to malware (#7) represents loss of control before any data breach occurs. In contrast, successful SQL injection (#2) can immediately result in Loss of Confidentiality.

Reference: §6.3 (Central Event), §6.3.1 (The Consequence Chain)

See also: System Risk Event (SRE), Data Risk Event (DRE), Business Risk Event (BRE)

Loss of Integrity (LoI)

A Data Risk Event outcome where the standing of a record as a faithful and attributable representation has changed other than through its intended, authorised process — typically an attacker's unauthorized modification. From the attacker's perspective: "Data modified". This refined definition focuses on the outcome rather than the mechanism (e.g., "tampering" is a mechanism that produces the LoI outcome). DRE notation: **I** (general), refined into **Ii** (Incorrect State) or **If** (Misattributed State) when the failing property is evidenced on the record. Integrity is not Accuracy: a record that is wrong but was never altered outside its authorised process is a data-quality matter, not a DRE.

Loss of Integrity — Incorrect State (LoIi)

A refinement of Loss of Integrity where **correspondence or completeness** fails: the record no longer corresponds to what it represents, or is no longer complete. Read off the record: the content is wrong. Examples: a wiped or garbled field, a corrupted database, a value keyed wrongly on purpose by an insider, a transcript showing a command that never ran. DRE notation: **Ii** (refined) or **I** (general). Answered by validation, reconciliation and four-eyes controls. Not to be confused with a Misattributed State (**If**), where the content may be perfectly accurate but its provenance is false. The two are told apart by inspecting the record, never by who caused it.

Loss of Integrity — Misattributed State (LoIf)

A refinement of Loss of Integrity where **provenance or attribution** fails: the record claims an origin, author or authority it does not have, while its content may be perfectly accurate. Examples: an entry posted under someone else's open session, a forged log line, a document carrying a signature or timestamp it was not given, a tool-call transcript written by attacker code in the harness's name. DRE notation: **If** (refined) or **I** (general). Answered by segregation of duties, provenance logging and behavioural monitoring. This is the joint Parker's *Authenticity* names from the other side; TLCTC

derives it from the stopping rule rather than adding a category. Not to be confused with an Incorrect State (Ii).

M

Malicious Code

Code written with harmful intent, distinguished in TLCTC between:

- **Exploit Code:** Targets specific vulnerabilities to modify software behavior (#2/#3)
- **Malware Code:** Operates within expected execution paths for harmful purposes (#7)
- **Malware Software:** Comprehensive suite of tools (foreign code) that may incorporate multiple techniques, including exploit capabilities

Malvertising (*Industry Term*)

The use of online advertising to distribute malware or redirect users to malicious websites. In TLCTC: malvertising is a **delivery vector**, not a distinct threat category. It can deploy either exploits or malware depending on the attacker's strategy:

- If it exploits a browser/plugin vulnerability: #3 Exploiting Client → #7 Malware
- If it delivers malware directly (e.g., fake download): #9 Social Engineering → #7 Malware
- If it redirects to a credential-harvesting site: #9 Social Engineering → #4 Identity Theft

Reference: V1.9.1 Clarifications, Buzz-Word Refinement (#3)

See also: Drive-By Download, Watering Hole Attack, Phishing

Malware (#7)

A threat cluster where an attacker abuses the inherent ability of a software environment to execute foreign executable content, including inherently malicious Malware Code or legitimate tools/scripts when they execute attacker-controlled or otherwise foreign code ("dual-use"). The generic vulnerability is the software environment's designed capability to execute potentially untrusted 'foreign' code, scripts, or binaries. Distinguished from #2/#3 which use Exploit Code targeting implementation flaws, and from #1 which manipulates existing functions without executing foreign code/scripts/binaries.

Related reading: [CVE-2026-21510: Windows Shell SmartScreen bypass](#), [CrowdStrike 2025 Threat Hunting Report — TLCTC](#), [CrowdStrike 2025 Global Threat Report — TLCTC](#), [Cobalt Strike capabilities × TLCTC V2.0](#), [The File Type Fallacy — extension blocklists](#), [GovCERT-CH blocked filetypes × TLCTC](#)

Man in the Middle (#5)

A threat cluster where an attacker intercepts, eavesdrops on, modifies, or relays communication between two parties without their knowledge or consent, by exploiting a privileged position on the communication path. The generic vulnerability is the lack of sufficient control over the communication path, including the implicit trust placed in local networks and intermediate network

infrastructure in standard IP networking. Position might be gained locally (shared Wi-Fi) or by leveraging control over existing network intermediaries.

Related reading: [tlctc-mitm-encryption-scion.html](https://tlctc.mitm-encryption-scion.html)

Medium Velocity Class (V2.0)

A velocity classification where attack progression occurs within hours. Typical threat clusters: #9 (Phishing), #4 (Manual Credential Abuse). Control strategy utilizes SIEM alerting and analyst triage, as human response times are sufficient for detection and initial response. Example: Phishing campaigns followed by manual reconnaissance and lateral movement.

Mitigating Controls

In the Bow-Tie model: barriers on the right (effect) side that detect, contain, reduce impact, or enable recovery after the central event occurs. Corresponds to NIST CSF functions: RESPOND, RECOVER.

Reference: §6.1 (Bow-Tie Structure)

MITRE ATT&CK

A globally-accessible knowledge base of adversary tactics and techniques based on real-world observations. In the TLCTC framework, MITRE techniques are considered operational-level detail that map to the strategic-level threat clusters. TLCTC proposes enhancement through adding cluster mappings and typical velocity attributes to techniques.

Related reading: [MITRE ATT&CK & STIX integration with TLCTC](#), [TLCTC × MITRE ATT&CK — interactive explorer](#), [MITRE ATT&CK & TLCTC — detection meets risk](#), [TLCTC × CWE — interactive explorer](#), [TLCTC × CWE — context is king](#), [MITRE ATLAS × TLCTC integration](#), [MITRE ATT&CK for ML \(AML\) × TLCTC](#), [NIST AI RMF + MITRE ATLAS × TLCTC](#), [MITRE ATT&CK & STIX × TLCTC V2.0 — implementation guide](#), [CKC + ATT&CK + TLCTC — Holy Trinity of Defense](#), [D3FEND × TLCTC — the missing axis](#), [Grok AI 4.1 — independent TLCTC validation](#), [Cobalt Strike capabilities × TLCTC V2.0](#), [TLCTC v2.5 monster prompt — SOC & Detection](#)

MFA Bombing / MFA Fatigue (*Industry Term*)

An authentication bypass technique where an attacker, having obtained valid credentials, repeatedly triggers MFA push notifications to overwhelm the user into accidentally approving one. In TLCTC, this is a four-step sequence: #4 → #1 → #9 → #4:

1. **#4 Identity Theft:** Attacker uses stolen userID/password
2. **#1 Abuse of Functions:** Repeatedly triggering legitimate MFA challenge requests (abusing intended functionality, not a code flaw)
3. **#9 Social Engineering:** Psychologically manipulating the user through fatigue/annoyance into approving a request
4. **#4 Identity Theft:** Successfully obtaining and using the MFA token to complete authentication

This example demonstrates how TLCTC decomposes a single "buzzword attack" into its constituent generic vulnerabilities.

Reference: V1.9.1 §Attack Path Notation (MFA Bombing Example)

See also: Identity Theft (#4), Abuse of Functions (#1), Social Engineering (#9)

MitM Position

A controlled point on a communication path that enables interception, observation, modification, injection, replay, or protocol downgrade/stripping. The attacker has achieved the ability to influence communication between two endpoints.

Reference: §4.2.2 (Global Definitions), R-MITM (§4.2.5)

N

NIST CSF (Cybersecurity Framework)

The National Institute of Standards and Technology Cybersecurity Framework providing guidelines for managing cybersecurity risk. The TLCTC framework integrates with NIST CSF by mapping the 10 threat clusters to the five core functions (Identify, Protect, Detect, Respond, Recover) and the GOVERN function in CSF 2.0. TLCTC proposes formal adoption of the 10 clusters as the standard taxonomy for Threat Identification in the ID.RA (Risk Assessment) category.

Related reading: [NIST CSF × TLCTC — critical analysis](#), [NIST CSF × TLCTC — bridging the gap](#), [NIST NICE × TLCTC — workforce capabilities](#), [NIST SP 800-218 \(SSDF\) × TLCTC](#), [NIST threat definitions — structural gap](#), [NIST AI RMF + MITRE ATLAS × TLCTC](#), [22 NIST definitions of "threat" — TLCTC](#), [SABSA × TLCTC — architecture × threat ontology](#), [TLCTC vs 30+ standards & regulations](#), [TLCTC Control Matrix Manager](#), [Control Matrix for SME & Private \(TLCTC v2.3\)](#), [Capability-based planning via 10×\(6×2\) matrix](#), [TLCTC v2.5 monster prompt — CISO & Risk](#), [The Risk Appetite Conversation Your Board Isn't Having](#), [Strategic risk management implementation guide v2.1](#)

NIS2 (Network and Information Security Directive 2) (*Industry Term*)

The EU directive establishing cybersecurity risk management and incident reporting obligations for essential and important entities. In TLCTC: NIS2 is an **incident-triggered regulation** — its reporting obligations activate at E1 (Significant Incident / System Compromise) regardless of whether personal data is involved. This contrasts with GDPR, which is **data-triggered** (activates at E2 when PII is affected). NIS2 Art. 23 requires a 24-hour early warning, creating a shorter event chain length (E1→E3b = 2 events) compared to GDPR's 72-hour notification timeline (E1→E2→E3a = 3 events).

Reference: V1.9.1 §Cyber Threat Radars

Related reading: [CVE-2020-17103 — patch closed an effect, not a cluster](#), [NIS2 Directive × TLCTC — implementation guide](#), [NIS2 — TLCTC pain points & fixes](#), [EU cyber regulation needs a common taxonomy](#), [EU regulation \(NIS2/DORA/CRA\) vs TLCTC](#), [GDPR vs NIS2 — different trigger points](#), [TLCTC vs 30+ standards & regulations](#), [Logical impossibility of control-first regulation](#), [The Commit Is the CVE — silent fixes & the patch-gap collapse](#), [TLCTC v2.5 monster prompt — Regulators & Standards](#)

See also: DORA, Regulatory Trigger Point, Event Chain Length, E_n Event Notation

Normative Keywords

The keywords MUST, MUST NOT, REQUIRED, SHALL, SHALL NOT, SHOULD, SHOULD NOT, RECOMMENDED, MAY, and OPTIONAL in the TLCTC specification are interpreted as described in RFC 2119 / RFC 8174. When these keywords appear in lowercase, they carry their ordinary English meaning.

Reference: §4.2 (Normative Keywords)

Notation Systems

The TLCTC framework employs two complementary notation systems:

- **Strategic Notation:** Human-readable format using #X (e.g., #1, #10) for executive communication and risk assessment
- **Operational Notation:** Machine-readable format using TLCTC-XX.YY (e.g., TLCTC-01.00) for tool integration, automation, and SIEM

Both notations remain fully compatible and can be used interchangeably based on context.

O

Observed Δt

A Δt value computed from two concrete time observations.

Reference: §12.0.3

Operational Layer

The detailed implementation level where security controls are implemented, monitored, and adjusted. Includes specific vulnerability management, threat intelligence (using frameworks like MITRE ATT&CK), TTP mapping, attack path analysis, vulnerability management (CVE reports), incident response, security testing, and monitoring. Uses the machine-first naming convention TLCTC-XX.YY, where XX is the two-digit cluster number (01–10) and YY is the two-digit sub-cluster number (00–99), for tool integration, SIEM rules, automation, threat intelligence exchange, and detailed documentation.

Reference: §4.2.1 (Two-Layer Naming Convention)

Operational Risk (OpRisk) (*Industry Term*)

The broader category of risks arising from inadequate or failed internal processes, people, and systems, or from external events. In TLCTC: cyber risks are explicitly defined as a **subset** of operational risks. While cyber risk management focuses on threats from unauthorized or unknown entities (covered by the 10 TLCTC clusters), comprehensive risk management must also consider traditional IT risks (e.g., "software failure", "error in use", "abuse of rights" by authorized actors), compliance risks, and third-party risks. Since v2.5 the boundary is stated as the **cause-side partition** and enforced by **R-SCOPE**: failure or external event, Error in Use, and Abuse of Rights are operational-risk rows with no cluster; only the Attack row — an intended action outside any

conferred entitlement — is cyber risk and receives a cluster. Actions of entitled actors inside their grant are Abuse of Rights; the same actors reaching past their grant are attacks and classify like anyone else.

Reference: V1.9.1 §Introduction

Related reading: [The Adoboli Paradox — Cyber vs Operational Risk](#)

See also: Cyber Risk, Business Risk Event

Operational Security Layer

The layer of TLCTC that addresses specific vulnerabilities, techniques, and procedures. Contains concrete vulnerabilities (CVEs), operational techniques (TTPs), and indicators used in detection, response, and engineering. Corresponds to TLCTC-XX.YY where YY ≠ 00.

Reference: Axiom VIII (§2), §4.2.1

OAuth Attack (*Industry Term — Decomposition Required*)

A commonly used but imprecise term that conflates multiple distinct attack mechanisms targeting OAuth implementations. In TLCTC, the framework's precision requirement reveals that "OAuth attack" must be decomposed into its specific mechanism:

- **Misconfigured redirect_uri:** #1 Abuse of Functions (abusing legitimate OAuth functionality via configuration)
- **XSS stealing OAuth tokens:** #3 Exploiting Client → #4 Identity Theft
- **Phishing OAuth consent:** #9 Social Engineering → #4 Identity Theft
- **Authorization code injection:** #1 (if design abuse) or #2 (if server-side validation flaw)

This decomposition is not creating ambiguity — it is exposing the security industry's imprecise terminology and enabling more targeted control implementation.

Reference: V1.9.1 §F (Industry Term Decomposition)

See also: Identity Theft (#4), Abuse of Functions (#1), Exploiting Client (#3)

OWASP (Open Worldwide Application Security Project) (*Industry Term*)

A nonprofit foundation providing freely available resources for web application security, including the OWASP Top 10 list of critical web application security risks. In TLCTC: OWASP risks and testing methodologies are considered operational-level detail. The OWASP Top 10 categories map to TLCTC clusters — e.g., "Injection" maps to #2/#3, "Broken Authentication" maps to #4, "Security Misconfiguration" maps to #1. TLCTC notes that OWASP (like STRIDE) is "per se incomplete" and recommends always starting threat assessment with the 10 TLCTC clusters.

Reference: V1.9.1 §Operational Layer

See also: STRIDE, Sub-Threat, Operational Layer

P

Parallel Operator (+)

Notation that denotes **concurrent** (or effectively concurrent) steps—actions that occur in the same phase where their ordering is not meaningful. Parallel steps **MUST** be grouped using parentheses: `(#X + #Y)`. Each element inside a parallel group is a separate Attack Step mapped to exactly one cluster.

Reference: §4.2.2 (Global Definitions), §11.2 (Parallel Operator)

Parallel Steps

Two or more clusters occurring simultaneously or in tight coordination within the same attack phase. Use when distinct generic vulnerabilities are exploited concurrently rather than sequentially.

Reference: §4.2.2 (Global Definitions), §11.2 (Parallel Operator)

Pass-the-Hash / Pass-the-Ticket (*Industry Term*)

Attack techniques where an attacker uses captured NTLM hashes (Pass-the-Hash) or Kerberos tickets (Pass-the-Ticket) to authenticate as a legitimate user without knowing the actual password. In TLCTC: the **acquisition** of the hash/ticket maps to the enabling cluster (e.g., #7 if extracted by malware, #1 if via lsass dump using a legitimate tool). The **use** of the hash/ticket to authenticate always maps to #4 Identity Theft per R-CRED.

Reference: V1.9.1 Buzz-Word Refinement (#4)

See also: Identity Theft (#4), Credential Acquisition, Credential Application

Password Spraying (*Industry Term*)

An attack that tries a small number of commonly used passwords against many accounts simultaneously, avoiding account lockout thresholds. In TLCTC: maps to #4 Identity Theft — the attacker is attempting to derive valid credentials to impersonate a legitimate identity. The generic vulnerability is the insufficient binding, at the point of authentication, between a presented credential and the authentic holder: a guessed valid password is accepted as if from the legitimate holder when controls such as MFA or lockout are absent (predictable passwords widen the guess space).

Reference: V1.9.1 Buzz-Word Refinement (#4)

See also: Identity Theft (#4), Brute-Force Attack

Patient Zero (*Industry Term*)

In TLCTC context: the first system compromised in an attack, representing the initial entry point before lateral movement occurs. Patient Zero systems are critical because they are typically exposed systems that cannot be fully protected by Umbrella Controls alone and require robust Local Controls. After Patient Zero compromise, attacks follow the lateral movement paradigm through the internal network.

Reference: V1.9.1 §A (Umbrella Controls)

See also: Local Controls, Umbrella Controls, Lateral Movement

Phishing (*Industry Term*)

A social engineering technique using deceptive communications (email, SMS, voice) to trick individuals into taking actions that compromise security. In TLCTC: phishing is the **delivery vector**, not a distinct threat category. The phishing lure/deception phase maps to #9 Social Engineering. Subsequent steps map to their respective clusters depending on the payload:

- Credential harvesting: #9 → #4 (phishing → identity theft)
- Malware delivery: #9 → #7 (phishing → malware execution)
- Browser exploit: #9 → #3 (phishing → client exploitation)
- Feature misconfiguration: #9 → #1 (phishing → abuse of functions)

Variants include Spear Phishing, Whaling, Vishing, and Smishing.

Reference: V1.9.1 Clarifications, Buzz-Word Refinement (#9)

See also: Social Engineering (#9), Spear Phishing, Vishing, Smishing

Physical Attack (#8)

A threat cluster where an attacker gains unauthorized physical interaction with or causes physical interference to hardware, devices, facilities, or data transmission media (including wireless signals). The generic vulnerability is the physical accessibility of infrastructure and the exploitability of physical-layer properties. Encompasses two sub-cluster vectors:

- **Mechanical vector (#8.1, contact):** Requires physical touch or direct interaction (tampering, theft, physical intrusion, unauthorized device connection)
- **Signal vector (#8.2, no contact):** Exploits physical properties without direct contact (TEMPEST, signal jamming, acoustic attacks, environmental disruption)

Related reading: [CWE-514 \(covert channel\) → TLCTC #8](#), [Topology of cyber attacks — Bridge vs Internal](#)

Pineapple Attack (*Industry Term*)

An attack using a Wi-Fi Pineapple (or similar rogue access point device) to create fake wireless networks that victims connect to, enabling traffic interception. In TLCTC: the physical deployment of the device maps to #8 Physical Attack (requiring physical proximity). Once victims connect, the attacker gains a MitM position — exploitation of that position maps to #5 Man in the Middle. The full sequence is #8 → #5. Some variants also involve #4 (faking SSID identity).

Reference: V1.9.1 Buzz-Word Refinement (#5)

See also: Physical Attack (#8), Man in the Middle (#5), Rogue Hotspot

Ping of Death (*Industry Term*)

A denial-of-service attack that sends malformed or oversized ICMP packets to crash or destabilize a target system. In TLCTC: if the crash results from an implementation flaw (buffer overflow in ICMP

handling), maps to #2 Exploiting Server or #3 Exploiting Client per R-ROLE and R-FLOOD. If the primary mechanism is volume-based, maps to #6 Flooding Attack.

Reference: V1.9.1 Buzz-Word Refinement (#2)

See also: Flooding Attack (#6), Implementation Defect (Availability Context)

Position Acquisition vs Position Exploitation

For #5 Man in the Middle: **Gaining** a MitM position maps to another cluster (#1, #8, #9, #10, or #2/#3 depending on initial generic vulnerability). **Exploiting** a MitM position (intercept, modify, relay, inject, replay, downgrade actions) maps to #5.

Reference: §4.2.2 (Global Definitions), R-MITM (§4.2.5)

Preventive Controls

In the Bow-Tie model: barriers on the left (cause) side that reduce likelihood of threats reaching the central event. Corresponds to NIST CSF functions: IDENTIFY, PROTECT.

Reference: §6.1 (Bow-Tie Structure)

Pretexting (*Industry Term*)

A social engineering technique where an attacker creates a fabricated scenario (pretext) to manipulate a target into divulging information or performing actions. In TLCTC: maps to #9 Social Engineering — the generic vulnerability is human psychological factors (trust, authority bias). Pretexting is a sub-threat of #9; any subsequent technical steps map to their own clusters.

Reference: V1.9.1 Buzz-Word Refinement (#9)

See also: Social Engineering (#9), Phishing

Privilege Escalation (*Industry Term*)

A commonly used but multi-faceted term describing an attacker gaining higher-level permissions than initially authorized. In TLCTC: privilege escalation is **not a single cluster** — it maps to different clusters depending on the technique employed:

- **Via software vulnerability** (buffer overflow, injection flaw): #2 Exploiting Server or #3 Exploiting Client per R-ROLE
- **Via misuse of legitimate features** (misconfiguration, overly permissive defaults): #1 Abuse of Functions
- **Via stolen credentials** (using admin credentials): #4 Identity Theft
- **Via social engineering** (manipulating users into granting access): #9 Social Engineering

Distinguishing the underlying technique allows for more targeted control implementation. The Intra-System Boundary Operator (`[[privilege][@from→@to]]`) can annotate privilege escalation in attack paths without changing cluster classification.

Reference: V1.9.1 Clarifications

Related reading: [AD → Domain Admin → Ransomware cascade, Calif M5: #2 → #2 \(Kernel's Role pt 1\)](#), [CVE-2026-31431 \(Copy Fail\): Linux kernel AF ALG](#), [CVE-2026-46300 \(Fragnesia\): Linux kernel XFRM](#)

See also: Abuse of Functions (#1), Exploiting Server (#2), Identity Theft (#4), Intra-System Boundary Operator

Process Injection (*Industry Term*)

A technique where an attacker inserts code into the address space of another running process. In TLCTC: process injection maps to different clusters depending on the mechanism:

- **Via designed features** (debugging APIs, DLL injection via legitimate Windows functionality): #1 Abuse of Functions — the injection capability was intentionally designed.
- **Via implementation flaws** (buffer overflows enabling code injection): #2 Exploiting Server or #3 Exploiting Client — the injection was never intended.

The key distinction is whether the injection vector was a designed feature being misused versus an underlying software vulnerability being exploited.

Reference: V1.9.1 Clarifications

See also: Abuse of Functions (#1), Exploiting Server (#2), Implementation Flaw

Programmer

A development role focused on architecture and strategy, responsible for designing overall software architecture and component interactions, making strategic decisions about frameworks and protocols, establishing secure coding standards and security requirements, and considering system-wide security implications. Primary responsibility for addressing threat clusters #1, #4, #5, #10 at an architectural level. Contrasts with the Coder role which focuses on implementation and craftsmanship.

Related reading: [Taming SonarQube with TLCTC \(Secure-Dev pt 2\)](#), [Programmer vs Coder in TLCTC \(Secure-Dev pt 1\)](#), [Adding the Developer's View to TLCTC](#)

Propagated PR (V2.0)

A Protection Requirement that "propagates backward" from a downstream event into the RS (Respond) container of an earlier event due to regulatory or policy requirements. Notation: $RS(E_n) = \{ \text{Response} \} \cup \{ \text{Propagated PR}(E_{n+1}) \} \cup \{ \text{Propagated PR}(E_{n+x}) \}$.

This mechanism explains how multiple regulatory obligations stack into a single incident response workflow. Example: A ransomware attack triggers E1 (System Compromise). If PII is affected, E2 (Data Risk Event) occurs, propagating GDPR notification requirements back into RS(E1). Simultaneously, if the organization is NIS2-scoped, the incident itself propagates NIS2 reporting into RS(E1). The result: two separate Propagated PR controls in the same RS container, with different timelines (72h vs 24h+72h) and different authorities. See also: RS Container, E_n Event Notation.

Related reading: [Propagated Controls — Rule of Propagation](#)

Protection Ring Architecture

The layered privilege model in computing systems (Ring 0 through Ring 3) where each ring represents a different privilege level. TLCTC framework analyzes threats at ring boundary interactions: Ring 0 (Kernel Mode), Ring 1 (HAL/Driver Level), Ring 2 (OS Services), Ring 3 (User Mode). Nine threat clusters (excluding #9 Social Engineering) apply at each boundary, with roles (client/server) determined by the direction of interaction across rings.

R

R-ABUSE (Function Misuse Determination) (Deprecated alias)

Retired v2.0 whitepaper rule ID; not part of the v2.5 normative registry. Its substance is carried unchanged by the #1 cluster definition and boundary tests (core paper §4). The ID keeps this original meaning and is never reused.

Original statement: If the attacker's success does not require any implementation flaw and instead abuses intended functionality, scope, or configuration via standard interfaces using expected input types, the step **MUST** be classified as #1 Abuse of Functions.

Reference: whitepaper §4.2.5 (R-ABUSE); superseded by core paper §4 (#1 boundary tests)

R-CRED (Credential Lifecycle Non-Overlap)

Global mapping rule: Credential acquisition maps to the enabling cluster; credential application **MUST** always map to #4 Identity Theft, **provided the identity claimed is not the presenter's own**. If both occur, they **MUST** be represented as at least two steps: (enabling cluster) → #4.

Self-issued identity (R-CRED proviso). A credential issued to the presenter by the target system through a designed enrolment function makes the presenter its authentic holder; using it is authentication as self and is **NOT** #4. Where the enrolment function granted the identity or its permissions outside their intended population or scope, that enrolment step maps to #1. Fictitious or pseudonymous self-registration is #1 (no identity impersonated); enrolment completed AS an existing identity is #1 → #4. The higher-abstraction test: is the system *deceived about who is authenticating*? Deceived → #4; not deceived (it enrolled this principal itself) → the vulnerability is elsewhere, usually #1.

Reference: §4.2.5 (R-CRED)

Related reading: [Cobalt Strike capabilities × TLCTC V2.0](#)

R-EXEC (Foreign Execution Recording Rule)

Global mapping rule: Whenever Foreign Executable Content (FEC) is interpreted, loaded, or executed, a #7 Malware step **MUST** be recorded at the moment of execution, independent of how execution was enabled. #7 is additive (does not replace the enabling cluster).

Reference: §4.2.5 (R-EXEC)

Related reading: [Cobalt Strike capabilities × TLCTC V2.0](#), [The File Type Fallacy — extension blocklists](#), [GovCERT-CH blocked filetypes × TLCTC](#)

R-FLOOD (Capacity Exhaustion vs Implementation Defect)

Global mapping rule: If the primary mechanism is volume or intensity exhausting finite resources, classify as #6 Flooding Attack. If the primary mechanism is an implementation defect that causes crash/hang/degradation (including algorithmic complexity), classify as #2 or #3 per R-ROLE.

Reference: §4.2.5 (R-FLOOD)

R-CHANNEL (Channel Control vs Code Flaw)

Global mapping rule (v2.5): If the defective logic is itself a communication-path control — peer authenticity (certificate validation, chain of trust, hostname matching, expiry or revocation checking), channel encryption, or algorithm negotiation — the generic vulnerability is the lack of sufficient control over the communication path and the weakness classifies as #5 Man in the Middle, not as #2 or #3 under R-ROLE. R-ROLE governs only where the defect is incidental to the control rather than constitutive of it (for example, memory corruption in a TLS parser).

R-CHANNEL classifies the *weakness*; R-MITM sequences the *attack path* (position acquisition versus action). The two do not conflict.

Reference: §6.1 (R-CHANNEL)

R-SCOPE (Entitlement Scope Boundary) (v2.5)

Admission rule for the whole registry: a step is classified under a cluster only where the action fell **outside the entitlement envelope** an accountable grantor actually conferred for it. Cause-side classification asks three questions in strict order — is there an actor; did they intend the outcome; did an accountable grantor confer an entitlement covering *this* action — and only an intended, unentitled action enters the Attack row where the clusters apply (implementation flaw required → #2/#3 per R-ROLE; otherwise #1). No actor is failure or external event; an unintended outcome is Error in Use; an intended in-grant action is **Abuse of Rights** — operational risk, no cluster, no System Risk Event, the consequence chain starting at the DRE. Entitlement means entitled, not permitted, and attaches to the person, not the token: credential use by anyone other than the grantee is never inside an envelope (#4 per R-CRED, subsequent function use #1), and exploiting an implementation flaw is never inside any grant. Not an actor test (Axiom IV). Entitlement is tested after intent and before the code-flaw test because an entitled actor exploiting a code flaw is an attack.

Reference: Core paper §3.5, §6.1; dictionary `rules[R-SCOPE]`, `cause_side_partition`

See also: Cause-Side Partition, Abuse of Rights, Entitlement, Error in Use, R-CRED

R-SUBSTRATE (Physical Property vs Implemented Logic)

Global mapping rule (v2.5): Classify as #8 Physical Attack only where a physical-layer property of the substrate — charge, voltage, electromagnetic emission, temperature, emission-borne timing, wear, or material state — is itself the exploited generic vulnerability. Where the physical layer serves only as the readout channel for a defect in implemented logic, classify by that defect (#2 or #3 per R-ROLE). Attacker proximity or possession is **not** the test.

The discriminating question is whether the attack is against the *implemented logic* or against the *physical representation* that logic runs on. Rowhammer is #8 (charge migration between adjacent DRAM cells is the vulnerability; nothing logical fails) even though it can be mounted from JavaScript. Spectre is #2 (speculation crosses an isolation boundary the design was meant to enforce; cache timing is only the readout). Power side-channel analysis is #8, because the cryptography is correct and the emission itself is the vulnerability.

R-SUBSTRATE is the *admission* test — whether a weakness qualifies as #8 at all. The sequencing principle formerly stated as R-PHYSICAL (now a deprecated alias) still holds — a qualifying physical step is #8 and subsequent technical steps are classified separately. They are complementary.

Reference: §6.1 (R-SUBSTRATE)

R-HUMAN (Human Manipulation Isolation) (Deprecated alias)

Retired v2.0 whitepaper rule ID; not part of the v2.5 normative registry. Its substance is carried unchanged by the #9 cluster definition and boundary tests (core paper §4) together with Axiom VI. The ID keeps this original meaning and is never reused.

Original statement: If the attacker's advantage comes from psychological manipulation of a human, that manipulation step **MUST** be classified as #9 Social Engineering, and any subsequent technical steps **MUST** be classified separately.

Reference: whitepaper §4.2.5 (R-HUMAN); superseded by core paper §4 (#9 boundary tests)

R-INTRA (Intra-System Boundary Rules) (V2.1)

The v2.5 normative registry carries exactly **two** R-INTRA rules governing the intra-system operator (|...|):

Rule	Summary
R-INTRA-7	Intra-system boundary crossings never change cluster classification. They are observability annotations, not classification inputs.
R-INTRA-9	The <code>memory</code> intra-system boundary type is deferred and MUST NOT be used.

The remaining usage guidance from the v2.1 drafting is non-normative and preserved as notation practice: the operator applies only to boundaries within a single system instance, attaches to the cluster step that accomplishes the crossing, never appears standalone, is optional (mainly for forensic or vendor-facing precision), and multiple annotations may follow one step. If a crossing requires a **separately evidenced vulnerability**, a new cluster step **MUST** be added — this follows directly from Axiom VI (see also SG-6) and is not an R-INTRA rule.

Numbering erratum (do not cite the v2.1 draft numbers). Early v2.1 drafts circulated a nine-rule series R-INTRA-1...9 in which "no cluster change" was numbered R-INTRA-4 and "distinct vulnerabilities" was numbered **R-INTRA-7**. In the canonical registry (v2.3 onward) R-INTRA-7 denotes the *no-cluster-change* rule. The draft series is therefore withdrawn: citing "R-INTRA-7" in its draft meaning contradicts the canon. Only the canonical meanings above are valid; this is the single known case in TLCTC

history of a rule ID changing propositions between draft and canon, recorded here so it can never happen silently.

Reference: core paper §6.2 (R-INTRA-7, R-INTRA-9), §11.3.6 (Intra-System Boundary Operator)

R-MITM (Position vs Action)

Global mapping rule: The method of gaining a privileged communication-path position maps to another cluster. #5 Man in the Middle begins only once the attacker controls a point on the communication path and performs MitM actions.

Reference: §4.2.5 (R-MITM)

R-PHYSICAL (Physical Domain Isolation) (Deprecated alias)

Retired v2.0 whitepaper rule ID; not part of the v2.5 normative registry. Its sequencing substance is carried by the #8 cluster definition and boundary tests (core paper §4); its admission question is now settled normatively by R-SUBSTRATE. The ID keeps this original meaning and is never reused.

Original statement: If the attacker's advantage comes from unauthorized physical interaction or interference with hardware, facilities, media, or signals, that step **MUST** be classified as #8 Physical Attack, and subsequent technical steps **MUST** be classified separately. Note that R-SUBSTRATE corrects a latent misreading of this phrasing: attacker physical access or proximity is **NOT** required for #8.

Reference: whitepaper §4.2.5 (R-PHYSICAL); superseded by core paper §4 (#8 boundary tests) and R-SUBSTRATE (core §6.1)

R-ROLE (Server vs Client Determination)

Global mapping rule: If the vulnerable component accepts and handles inbound requests relative to the attacker, classify as #2 Exploiting Server. If the vulnerable component consumes external responses/content relative to the attacker, classify as #3 Exploiting Client. Roles are established by call direction at **any** interface, including intra-system privilege interfaces (syscall, hypercall, IPC, driver IOCTL) — a network is not a precondition. A kernel handling a crafted syscall from a lower-privileged process is in server-role (#2). Per R-INTRA-7, the boundary crossing itself remains an observability annotation and **MUST NOT** be treated as a classification input.

Reference: §4.2.5 (R-ROLE)

Related reading: [Calif M5: #2 → #2 \(Kernel's Role pt 1\)](#), [CVE-2025-21333 revisited: #2 after all \(Kernel's Role pt 2\)](#), [Apache 2.4.67 — 11 CVEs decomposed](#)

R-SUPPLY (Trust Acceptance Event Placement)

Global mapping rule: #10 Supply Chain Attack **MUST** be placed at the Trust Acceptance Event (TAE) —the moment where the third-party trust link is honored and the trust artifact becomes authoritative inside the organization's domain.

Reference: §4.2.5 (R-SUPPLY)

R-TRANSIT (Transit Boundary Rules) (V2.1)

The complete transit boundary rule set governing use of the transit operator (⇒):

Rule	Name	Summary
R-TRANSIT-3	Vendor Code on Target Device	Vendor code running on the target device is NOT transit — it is the attack surface and MUST be classified by R-ROLE

Only **R-TRANSIT-3** is part of the v2.5 normative registry. The remaining v2.1 drafting guidance is non-normative notation practice: a transit party must be distinct from source and target and actually sit between them in the delivery path; the annotation is optional (recommended where the intermediary has meaningful control responsibility, omissible for pure conduits); compromise or coercion of the intermediary must be modeled as its own preceding cluster step; transit annotations never change cluster classification (see SG-2); and chained transit may be used when each party has independent analytical relevance. The withdrawn draft IDs R-TRANSIT-1, -2, -4, -5, -6, -7, -8 are not to be cited.

Reference: core paper §6.2 (R-TRANSIT-3), §11.3.5 (Transit Boundary Operator)

Realtime Velocity Class (V2.0)

A velocity classification where attack progression occurs within seconds or milliseconds. Typical threat clusters: #6 (Flooding), #2 (Wormable Exploits). Control strategy must rely on architecture, hardening, and circuit breakers, as detection and response times are insufficient. Example: DDoS attacks, wormable exploits like EternalBlue, automated exploitation frameworks.

Regulatory Trigger Point (V2.0)

The specific event type in a TLCTC event chain that activates a regulatory notification or compliance obligation. Different regulations have different trigger points within the same attack sequence:

- **Data-triggered regulations** (e.g., GDPR Art. 33): Activate at E2 (Data Risk Event involving PII) — no PII affected means no GDPR notification obligation
- **Incident-triggered regulations** (e.g., NIS2 Art. 23): Activate at E1 (Significant Incident / System Compromise) — regardless of whether PII is involved

Understanding regulatory trigger points enables CISOs to build precise IR playbooks mapping specific RS container actions to logical triggers rather than generic "reporting checklists". See also: Propagated PR, Event Chain Length.

Responsibility Sphere

The organizational owner of a domain, denoted as @Entity. Examples: @Org, @Vendor, @Facilities, @HR, @CloudProvider, @MSP. Different spheres have different policies, teams, governance structures, and potentially different legal boundaries. Domain boundary definitions identify where responsibility and control shift during an attack, which is critical for incident response, forensics, and legal responsibility. Defined in tlctc-responsibility-spheres.json and customizable per organization. Standard spheres include: Attacker Side, Third-Party/Vendor Side, Victim Side, Shared/Transit. Used in conjunction with the domain boundary operator (||) in attack path notation.

Reference: §4.2.2 (Global Definitions), §11.4, §5.1.2

Risk Event

In the TLCTC Bow-Tie model, the central occurrence that represents the materialization of a threat, positioned between causes (threats) and effects (consequences). For cyber risk, this is "Loss of Control" or "System Compromise". Can trigger event chains where one outcome becomes the event triggering subsequent events.

Role Determination

Classification of a component as server-role or client-role based on its behavior in the specific interaction being classified. The same software product MAY appear as server-role in one interaction and client-role in another. Classification MUST follow the role of the component being exploited in the step.

Reference: §4.2.2 (Global Definitions), R-ROLE (§4.2.5)

RS Container (Respond Container) (V2.0)

The logical collection of RESPOND-function controls and actions for a specific event (E_n) in the TLCTC event chain. An RS Container holds:

- **Direct Response Actions:** Containment, eradication, forensics for the event itself
- **Propagated PR Controls:** Protection requirements inherited from downstream compliance events (E3a, E3b, etc.)

Notation: $RS(E_n) = \{ \text{Response} \} \cup \{ \text{Propagated PR}(E_{n+1}) \} \cup \{ \text{Propagated PR}(E_{n+x}) \}$. Example: $RS(E_1)$ for a ransomware incident may contain both incident containment actions AND propagated GDPR/NIS2 notification controls, each with distinct timelines and reporting authorities. Aligns with NIST CSF RESPOND function. See also: Propagated PR, Regulatory Trigger Point.

Ransomware (*Industry Term*)

Malware that encrypts a victim's data and demands payment for the decryption key. In TLCTC: the execution of the ransomware binary maps to #7 Malware (foreign code executed via designed execution capability). The resulting data encryption is a Data Risk Event: [DRE: Ac] (Loss of Accessibility — data exists but is unusable). Note: this is Loss of **Accessibility**, not Loss of Availability (the encrypted files still exist on disk). A full ransomware attack path typically involves multiple clusters, e.g., #9 → #7 → #7 → #4 → #1 → #7 as seen in the Emotet/Ryuk case study (the final #1 → #7 is sequential — function abuse deploys the payload, which then executes; per §11.2.2 a known order uses →, not a parallel group).

Reference: V1.9.1 §Definitions (#7), §Data Risk Event Types, §E (Emotet@Heise)

Related reading: [AD → Domain Admin → Ransomware cascade](#), [Chaos Ransomware — TLCTC forensic](#), [Verizon DBIR 2025 — TLCTC](#), [20 annotated attack paths \(Ransomware, BEC, OT, ...\)](#)

See also: Malware (#7), Accessibility (Data Risk Event), Loss of Accessibility

RCE (Remote Code Execution) (*Industry Term — Decomposition Required*)

A commonly used but imprecise term describing CVEs that enable an attacker to execute arbitrary code on a remote target. In TLCTC, "RCE" always conflates the vulnerability with its exploitation and MUST be decomposed:

- The **vulnerability** that enables execution: #2 Exploiting Server (server-side flaw) or #3 Exploiting Client (client-side flaw)
- The **actual execution** of foreign code: #7 Malware
- Correct notation: #2 → #7 or #3 → #7

This decomposition is essential because it identifies two distinct generic vulnerabilities being exploited, each requiring different controls.

Reference: V1.9.1 §F (Industry Term Decomposition)

Related reading: [Calif M5: #2 → #2 \(Kernel's Role pt 1\)](#), [CVE-2025-21333 revisited: #2 after all \(Kernel's Role pt 2\)](#), [Apache 2.4.67 — 11 CVEs decomposed](#)

See also: Exploiting Server (#2), Exploiting Client (#3), Malware (#7)

RFID Skimming (*Industry Term*)

The unauthorized reading of RFID (Radio-Frequency Identification) chips from proximity, typically to clone access cards or extract stored data. In TLCTC: maps to #8 Physical Attack (specifically the #8.2 signal vector, no contact) — the attacker exploits physical properties (radio frequency emanations) without direct contact with the target device.

Reference: V1.9.1 Buzz-Word Refinement (#8)

See also: Physical Attack (#8), TEMPEST

Risk Appetite / Risk Tolerance (*Industry Term*)

Risk Appetite: The level and type of cyber risk an organization is willing to accept in pursuit of its objectives. **Risk Tolerance:** The acceptable deviation from the defined risk appetite. In TLCTC: risk appetite and tolerance should be defined per threat cluster, enabling differentiated risk treatment. For example, an organization may have zero tolerance for #10 Supply Chain risks in critical infrastructure but moderate tolerance for #6 Flooding risks with adequate DDoS mitigation in place. Risk appetite is a Strategic Management Layer concern that informs operational priorities.

Reference: V1.9.1 §Strategic Management Layer

See also: Strategic Management Layer, GOVERN (GV)

Rogue Hotspot (*Industry Term*)

A fraudulent Wi-Fi access point set up by an attacker to intercept traffic from unsuspecting users. In TLCTC: the physical deployment of the rogue access point maps to #8 Physical Attack (physical proximity required). Interception of traffic through the rogue AP maps to #5 Man in the Middle. If the attacker fakes an SSID to impersonate a legitimate network, that may additionally involve #4 (identity impersonation). Typical sequence: #8 → #5.

Reference: V1.9.1 Buzz-Word Refinement (#5)

See also: Pineapple Attack, Man in the Middle (#5), Physical Attack (#8)

Rootkit (*Industry Term*)

Malware designed to provide continued privileged access to a system while actively hiding its presence. In TLCTC: maps to #7 Malware — the rootkit is foreign executable content running via the system's designed execution capability. The persistence and stealth mechanisms are operational details within the #7 cluster.

Reference: V1.9.1 Buzz-Word Refinement (#7)

See also: Malware (#7), Spyware

S

SAST (Static Application Security Testing) (*Industry Term*)

A testing methodology that analyzes application source code, bytecode, or binary code for security vulnerabilities without executing the program. In TLCTC: SAST is a **preventive control** (IDENTIFY/PROTECT) primarily targeting #2 Exploiting Server and #3 Exploiting Client by identifying implementation flaws before deployment.

See also: DAST, Control, Exploiting Server (#2), Exploiting Client (#3)

SBOM (Software Bill of Materials) (*Industry Term*)

A formal, machine-readable inventory of all software components, libraries, and dependencies used in a software product. In TLCTC: SBOM is a critical control for #10 Supply Chain Attack — it enables organizations to identify and track third-party components, verify their integrity, and rapidly assess exposure when vulnerabilities are discovered in dependencies.

See also: Supply Chain Attack (#10), SCA, Third-Party Trust Link

SCA (Software Composition Analysis) (*Industry Term*)

Automated tools that identify open-source and third-party components in a codebase, flagging known vulnerabilities and license compliance issues. In TLCTC: SCA is a preventive control for #10 Supply Chain Attack, enabling detection of vulnerable or malicious dependencies before they are integrated into production systems.

See also: Supply Chain Attack (#10), SBOM

Semantic Guardrails (SG-1 through SG-7) (V2.1)

Normative rules that prevent V2.1 operators (transit boundary, intra-system boundary, unresolved-step) from drifting the classification model. They enforce that: classification is cause-first (SG-1), topology is not classification (SG-2), annotations are subordinate to clusters (SG-3), effects are not threats (SG-4), actors are not threats (SG-5), distinct exploits require distinct steps (SG-6), and stripping V2.1 annotations must leave a valid cluster sequence (SG-7).

Reference: §4.2.4 (Semantic Guardrails)

Scope of Server Software

In TLCTC: includes Server APIs, incorporated Library APIs, Socket APIs, and Local APIs that run on server-side systems to provide services and resources to clients. The scope defines the attack surface for #2 Exploiting Server.

Reference: V1.9.1 Clarifications

See also: Exploiting Server (#2), Scope of Client Software

Scope of Client Software

In TLCTC: encompasses Client APIs, incorporated Library APIs, Socket APIs, and Local APIs that operate on the client side of a communication. The scope defines the attack surface for #3 Exploiting Client. Client software includes web browsers, mobile apps, desktop applications, document readers, command-line clients, client libraries, API consumers, and background services acting as clients.

Reference: V1.9.1 Clarifications

See also: Exploiting Client (#3), Scope of Server Software

Secure Software Development Life Cycle (SSDLC)

A structured approach to embedding security throughout the software development process. The TLCTC framework integrates into each SSDLC phase, with programmer-level decisions during Requirements and Design, coder-level implementation during the Implementation phase, and both roles contributing to verification during Testing and ongoing vigilance during Maintenance.

Related reading: [Taming SonarQube with TLCTC \(Secure-Dev pt 2\)](#), [Programmer vs Coder in TLCTC \(Secure-Dev pt 1\)](#), [SSDLC for developers – the "S" problem](#), [SSDLC phase-by-phase reference – TLCTC v2.1](#), [Adding the Developer's View to TLCTC](#)

Sequence

The ordered progression of threat clusters in an attack. The TLCTC framework recognizes that identified Top-Level Threats must also be seen as sequence components in the attack scenario, with attackers using these components in varying orders depending on their script (Axiom VIII). Example: A phishing attack might follow #9→#3→#7, or a more complex attack might be #9→#7→#4→(#1+#7).

Sequence Operator (→)

The operator meaning: the right-hand step occurs after the left-hand step, and the left-hand step enables or makes possible the right-hand step in the described scenario. ASCII alternative: ->.

Reference: §11.1 (Sequence Operator)

Server-Role Component

A component that **accepts and handles inbound requests or stimuli** relative to the attacker. The component is in "server role" for the specific interaction being classified.

Reference: §4.2.2 (Global Definitions), R-ROLE (§4.2.5)

Session Hijacking (*Industry Term*)

An attack where an adversary takes over an active session by stealing or predicting session tokens/cookies. In TLCTC: the **acquisition** of the session token maps to the enabling cluster (e.g., #5 if intercepted via MitM, #3 if stolen via XSS, #7 if captured by malware). The **use** of the stolen session to impersonate the legitimate user always maps to #4 Identity Theft per R-CRED.

Reference: V1.9.1 Buzz-Word Refinement (#4)

See also: Identity Theft (#4), Credential Acquisition, Credential Application

SIEM (Security Information and Event Management) (*Industry Term*)

A category of security tools that aggregate and analyze log data from across an organization's infrastructure to detect security events. In the TLCTC velocity model, SIEM is the primary detection tool for the **Medium Velocity Class** (hours) where analyst triage is feasible. SIEM uses operational notation (TLCTC-XX.YY) for correlation rules, while dashboards may display strategic notation (#x) for SOC managers.

See also: Medium Velocity Class, EDR, SOAR, Notation Systems

Slowloris (*Industry Term*)

An application-layer denial of service attack that holds many connections to the target web server open by sending partial HTTP requests, slowly exhausting server connection resources. In TLCTC: maps to #6 Flooding Attack — the primary mechanism is exhausting finite connection/thread pool capacity. Although each individual request is small, the aggregate effect overwhelms the server's connection handling capacity.

Reference: V1.9.1 Buzz-Word Refinement (#6)

See also: Flooding Attack (#6), DDoS, HTTP Flood

Smishing (*Industry Term*)

Social engineering attacks delivered via SMS text messages. In TLCTC: maps to #9 Social Engineering — the delivery channel (SMS) does not change the cluster classification. Subsequent steps map to their respective clusters (e.g., #9 → #4 for credential harvesting, #9 → #7 for malware delivery).

Reference: V1.9.1 Buzz-Word Refinement (#9)

See also: Social Engineering (#9), Phishing, Vishing

SOAR (Security Orchestration, Automation, and Response) (*Industry Term*)

A category of security tools that enable automated incident response through predefined playbooks. In the TLCTC velocity model, SOAR is critical for the **Fast Velocity Class** (minutes) where automated containment is necessary because human analyst response times are insufficient. SOAR playbooks can be structured around TLCTC attack paths to automate response actions specific to detected cluster sequences.

See also: Fast Velocity Class, EDR, SIEM

System Failure (v2.5)

The **Failure type** of the System Risk Event: loss of function with no actor holding capability — software or hardware failure, misconfiguration, capacity exhaustion without an attacker, an external event, or an unintended act (Error in Use) that breaks the system. It sits at the same altitude as System Compromise and passes through the same conditional gate ("only if data is affected") into the data layer, so it shares the consequence chain **SRE** → **DRE** → **BRE*** while having no cause-side classification: operational risk, no cluster, no attacker's view. Named in v2.5 so that the failure branch drawn at the same altitude in the framework's figures is inside the model rather than implicit.

Reference: Core paper §3.4; dictionary `system_risk_event.types[failure]`

See also: System Risk Event (SRE), Loss of Control / System Compromise, Error in Use, Cause-Side Partition

System Risk Event (SRE)

Any risk event at the **system altitude** — the point at which a system's behavior, privileges, data, or trust relationships depart from what its owner controls — and the central event in the TLCTC Cyber Bow-Tie model. Since v2.5 the SRE has **two types at one altitude: System Compromise / Loss of Control** (an actor holds capability over the system sufficient to pursue objectives; reached only through cluster steps; the pivot of the bow-tie) and **System Failure / Loss of Function** (no actor holds anything; operational risk, no cluster). Abuse of Rights produces neither type — the system was obeyed, not compromised, and did not fail — so its chain begins at the DRE. The SRE is the pivot point between the cause side (threat clusters exploiting generic vulnerabilities) and the consequence side (data and business risk events). It is the first event in the consequence chain **SRE** → **DRE** → **BRE***, where each transition has its own Δt representing a detection and intervention window. Not every SRE leads to a DRE — detection and containment at the central event can break the chain before data-level consequences materialize.

Disambiguation: "Loss of Control" is always abbreviated **SRE**, never "LoC". The abbreviation **LoC** is reserved exclusively for Loss of Confidentiality, a *consequence-side* Data Risk Event. See **Loss of Confidentiality (LoC)**.

Reference: §6.3 (Central Event), §6.3.1 (The Consequence Chain)

See also: Loss of Control / System Compromise, Data Risk Event (DRE), Business Risk Event (BRE)

System-Risk Layer (v2.5)

The layer of the Cyber Bow-Tie at which cluster steps are prevented, detected, and contained before a compromise matures into data-level and business-level consequences — the control-placement layer around the System Risk Event (SRE), between the cause side (threat clusters #1–#10) and the consequence chain (DRE → BRE*). Its baseline concept is the **TLCTC × NIST CSF control matrix**: ten clusters crossed with six CSF functions yield **at minimum 60 distinct control objectives**; organization-specific refinements may subdivide cells but never merge them below the 10 × 6 grid. This fixed floor enables two strategic integrations. **Metrics:** every cell is a named objective to which Key Control Indicators attach — technical KCIs measure the cell's state and

coverage, procedural KCIs (KPIs) its process performance — with targets derived from risk appetite. **Security architecture:** the umbrella-control column, read down all ten rows, defines the enterprise's shared security architecture and connects the matrix to secure development through the Programmer role, which sets the architectural frame and secure-coding standards under which the Coder implements.

The system-risk layer is a control-placement concept, not a taxonomy layer: it is distinct from the Strategic/Operational two-layer naming model (Axiom VIII) and does not alter any cluster definition or classification rule.

Reference: Application paper §8–§10, §15 (Part B/C)

See also: System Risk Event (SRE), Control Objective, Umbrella Controls, Local Controls, KCI (Key Control Indicator), Programmer

Social Engineering (#9)

A threat cluster where an attacker psychologically manipulates individuals into performing actions counter to their or their organization's best interests, such as divulging confidential information, granting access, executing code, or bypassing security procedures. The generic vulnerability is human psychological factors: gullibility, trust, ignorance, fear, urgency, authority bias, curiosity, or general compromisability. Often serves as the initial vector enabling other threat clusters (e.g., #9→#4 for credential harvesting, #9→#7 for malware installation, #9→#1 for feature misconfiguration).

Related reading: [CVE-2026-21510: Windows Shell SmartScreen bypass](#), [Topology of cyber attacks — Bridge vs Internal](#), [tlctc-social-engineering-bifurcation.html](#), [TLCTC+ for NCSCs & CERTs — national reporting](#)

Spear Phishing (Industry Term)

A targeted social engineering attack directed at specific individuals or organizations, using personalized information to increase credibility. Variants include **Whaling** (targeting senior executives). In TLCTC: maps to #9 Social Engineering — the targeting precision is an operational detail but does not change the cluster. Subsequent steps map to their respective clusters based on the payload delivered.

Reference: V1.9.1 Buzz-Word Refinement (#9)

See also: Social Engineering (#9), Phishing, Whaling

Spyware (Industry Term)

Malware that covertly monitors user activity, collects data (keystrokes, browsing history, credentials), and transmits it to the attacker. In TLCTC: maps to #7 Malware — foreign code executing via the system's designed execution capability. The data collection constitutes a Data Risk Event: [DRE: C] (Loss of Confidentiality). If collected credentials are subsequently used, an additional #4 Identity Theft step applies.

Reference: V1.9.1 Buzz-Word Refinement (#7)

See also: Malware (#7), Keylogger, Trojan

SQL Injection (*Industry Term*)

An implementation flaw where an attacker inserts malicious SQL statements into application queries through unvalidated input, enabling unauthorized database access. In TLCTC: maps to #2 **Exploiting Server** — a coding flaw in server-side query building that creates an unintended data→code transition. SQL injection can lead to immediate Data Risk Events: [DRE: C] (data exfiltration), [DRE: I] (data modification), or [DRE: Av] (data deletion).

Reference: V1.9.1 §Definitions (#2), Buzz-Word Refinement

See also: Exploiting Server (#2), Implementation Flaw, Command Injection

SSRF (Server-Side Request Forgery) (*Industry Term*)

An industry label for an effect, not for a generic vulnerability: a server is induced to issue requests to destinations the attacker chooses, such as cloud metadata endpoints, internal services, or RFC 1918 hosts. TLCTC therefore does **not** assign SSRF one cluster: it splits into #1 **Abuse of Functions** or #2 **Exploiting Server** depending on the abused capability. Decision rule: *did the abused capability execute as designed, just for a forbidden target?*

- **Yes** → #1 **Abuse of Functions**. The endpoint is meant to fetch a URL — image optimizer, link preview, webhook tester, document/PDF renderer, import-from-URL, a rewrites/proxy feature. The attacker supplies a valid input of an expected type, the fetch function runs correctly, and only the target is forbidden; inputs stay data and no foreign code is introduced (per the #1 boundary tests). A missing or bypassed allowlist is the *bypass primitive*, not the abused capability, and does not move the cluster. Examples: CVE-2026-44578 (Next.js WebSocket-upgrade SSRF), the Next.js image-optimizer SSRFs, Capital One 2019.
- **No** → #2 **Exploiting Server**. The outbound request exists only because server-side implementation is broken: URL-parser confusion, mishandled redirects, DNS rebinding in the fetch logic, or an XXE that makes the XML parser resolve an external entity. No designed fetch capability was used as intended — a flaw manufactured the request.

SSRF is never a complete threat statement (Axiom IV) — it is a position-acquisition step and must be written as a sequence. Cloud-credential theft via the instance metadata service is #1 → #4 (fetch function abused to read 169.254.169.254, returned IAM credentials then applied per R-CRED); reaching and then exploiting an unexposed internal service is #1 → #2; the XXE-driven variant starts at #2.

The split drives control placement: for #1 the fetch feature itself is the surface to defend (allowlist / `remotePatterns`, protocol and response-size limits, outbound egress policy, IMDSv2) — input validation and WAF signatures are weak there because the input is valid; for #2 the defence is fixing the flaw and hardening the trust check that failed to run.

Reference: Core paper §4 (#1 boundary tests); V1.9.1 Buzz-Word Refinement (#1, #2)

Related reading: [CVE-2026-44578: Next.js WebSocket SSRF](#)

See also: Abuse of Functions (#1), Exploiting Server (#2), XXE (XML External Entity) Injection, Implementation Flaw, Position Acquisition vs Position Exploitation

SSL Stripping (*Industry Term*)

A MitM technique where an attacker downgrades HTTPS connections to HTTP by intercepting and modifying communication between client and server. In TLCTC: SSL stripping is an action performed from a MitM position and maps to #5 Man in the Middle (exploiting the controlled communication path to perform a protocol downgrade). The attacker must already have the MitM position (gained via #1 ARP spoofing, #8 physical access, etc.), making the full sequence e.g., #1 → #5.

Reference: V1.9.1 Buzz-Word Refinement (#1)

See also: Man in the Middle (#5), ARP Spoofing, Position Acquisition vs Position Exploitation

STIX (Structured Threat Information Expression)

A standardized language for representing cyber threat information. TLCTC integrates with STIX at the **classification layer**: cluster identifiers can be carried as custom properties (e.g., `x_tlctc_primary_cluster`) or controlled-vocabulary tags on existing STIX objects such as `attack-pattern` and `indicator`. **Layer 3 attack path instances do not round-trip to STIX 2.1** — Δt annotations, parallel groups, boundary operators (`||...||`, `⇒`, `|...|`), DRE tags, and unresolved operators (`?`, `...`) have no native STIX 2.1 equivalents and cannot be reconstructed from `attack-pattern` + SRO decomposition. Full Layer 3 exchange requires a custom STIX extension or transporting the TLCTC JSON instance as an out-of-band artifact.

Related reading: [MITRE ATT&CK & STIX integration with TLCTC](#), [MITRE ATT&CK & STIX × TLCTC V2.0 — implementation guide](#)

STRIDE (*Industry Term*)

A threat modeling methodology developed by Microsoft that categorizes threats into six types: Spoofing, Tampering, Repudiation, Information Disclosure, Denial of Service, and Elevation of Privilege. In TLCTC: STRIDE is considered "per se incomplete" — it conflates causes with outcomes (e.g., "Information Disclosure" is an outcome/DRE, not a cause) and lacks coverage of human and physical threat vectors. TLCTC's cause-oriented 10-cluster model provides a more logically consistent foundation, and TLCTC recommends always starting threat assessment with its clusters rather than STRIDE alone.

Reference: V1.9.1 §Standardizing Strategic Cybersecurity, §Operational Layer

Related reading: [Beyond STRIDE — TLCTC superior approach](#), [Beyond STRIDE — upgrading Microsoft TM Tool](#), [End of semantic diffusion — DREAD vs STRIDE vs TLCTC](#)

See also: OWASP, TLCTC

SYN Flood (*Industry Term*)

A network-layer denial of service attack that exploits the TCP three-way handshake by sending many SYN requests without completing the handshake, exhausting the target's connection table. In TLCTC: maps to #6 Flooding Attack — volume exceeding finite connection capacity.

Reference: V1.9.1 Buzz-Word Refinement (#6)

See also: Flooding Attack (#6), DDoS, UDP Flood

Strategic Layer (Human-First)

A naming convention for TLCTC clusters using the format #x where $X \in \{1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$. Used for executive communication, risk registers, board reporting, strategic planning, and high-level attack path discussion.

Reference: §4.2.1 (Two-Layer Naming Convention)

Strategic Management Layer

The stable top-level layer of TLCTC containing the 10 clusters and their generic vulnerabilities. Used for governance, control mapping, and comparable incident documentation. Focuses on risk management, policy-making, and program governance, including threat cluster categorization, generic vulnerability identification, risk appetite and tolerance definition, security program management, compliance and governance, and resource allocation. Corresponds to cluster identifiers #1–#10 or TLCTC-XX.00. Uses strategic notation (#X) for executive communication.

Reference: Axiom VIII (§2), §4.2.1

Sub-Threat

Specific, detailed attack techniques or methods that fall within a broader Top Level Cyber Threat Cluster. Sub-threats represent the operational-level detail beneath the strategic-level clusters. Example: Under #2 Exploiting Server, sub-threats include SQL Injection, Buffer Overflows, RCE via Deserialization, Template Injection, and XXE Injection. (SSRF is deliberately absent: it is an effect label that splits between #1 and #2 — see SSRF.)

Related reading: [Dual-layer notation — TLCTC-XX.YY enumeration](#)

Supply Chain Attack (#10)

A top-level threat cluster on the cause side of the bow-tie, where an attacker compromises systems by abusing the trust relationship within an organization's supply chain. The attacker targets vulnerabilities in third-party software components, hardware, services, or distribution/update mechanisms that are **trusted and integrated** into the organization's own environment or products. The generic vulnerability is the necessary reliance on, and implicit trust placed in, external suppliers, vendors, components, and their associated development or distribution processes.

Supply Chain as "Bridge Not Bucket": #10 is a *bridge* threat cluster that marks the use of a trusted supply-chain channel as an attack vector to cross from one domain/trust boundary into another (e.g. @Vendor → @Org). It does *not* absorb the semantics of other clusters (#1–#9).

Three Key Supply-Chain Vectors (#10.x):

- **#10.1 Update Vector:** Post-deployment delivery/update flow compromise
- **#10.2 Development Vector:** Pre-deployment build/dev pipeline, repositories, or package ecosystem compromise
- **#10.3 Hardware Supply Chain Vector:** Hardware component, firmware, or manufacturing/assembly compromise

Control-Level Third-Party Dependencies (Not #10): Dependencies on third parties for patch delivery, security updates, or managed services are treated as *governance/control dependencies* on the right side of the bow-tie. They are **not themselves** a #10 Supply Chain Attack unless the trusted integration channel is directly abused as an attack vector.

Related reading: [Topology of cyber attacks — Bridge vs Internal](#), [20 annotated attack paths \(Ransomware, BEC, OT, ...\)](#), [blog-attack-path-supply-chain.html](#), [tlctc-npm-supply-chain.html](#)

System Compromise

Alternative term for "Loss of Control" in the Cyber Bow-Tie model. See **System Risk Event (SRE)** and **Loss of Control / System Compromise**.

T

Tailgating (*Industry Term*)

A physical social engineering technique where an unauthorized person follows an authorized person through a secured entrance. In TLCTC: maps to #9 Social Engineering when it relies on human psychology (politeness, trust) to gain access. If the entry is forced without social manipulation (e.g., slipping through a door), it maps to #8 Physical Attack. Subsequent technical exploitation maps to its own cluster.

Reference: V1.9.1 Buzz-Word Refinement (#9)

See also: Social Engineering (#9), Physical Attack (#8)

Tech Enablers Overlay (V2.1)

A forward-looking **informative overlay** on the Cyber Threat Radar that maps emerging technologies (e.g., agentic AI, quantum-resistant crypto, deepfake toolchains, commodity exploit kits) against two axes: the **cluster axis** (which generic vulnerability the technology amplifies, #1 – #10) and the **Actor Group axis** (Nation-State / Cybercriminal-Ransomware / Cybercriminal-General / Hactivist / Insider / Amateur-ScriptKiddie). Complements the Attacker Profile overlay: where profiles describe which clusters actors favor *today*, the Tech Enablers Overlay tracks which new capabilities are entering the ecosystem, for whom, and which clusters they are likely to amplify next. Entries carry adoption-level indicators (observed / emerging / hypothesized) and snapshot dates; they **MUST** map to clusters via the generic vulnerability amplified, not by vendor or product category. Shifts between snapshots **SHOULD** trigger a review of the organizational radar and relevant Attacker Profiles.

Reference: §17.4

See also: Cyber Threat Radar, Attacker Profile, Actor Group, Actor Archetype

Tool: </tools/tech-enablers-radar.html>

Techniques (TTPs)

Specific methods, procedures, and tactics that attackers use to exploit vulnerabilities and achieve their objectives. In the MITRE ATT&CK framework, techniques represent the operational "how" of attacks—the concrete actions adversaries take (e.g., T1190 "Exploit Public-Facing Application", T1566 "Phishing"). Techniques often reference specific vulnerabilities (CVEs) they exploit and the platforms/systems they target.

Relationship to TLCTC: Each MITRE technique can be mapped to one or more TLCTC clusters based on the generic vulnerability being exploited:

- T1190 (Exploit Public-Facing Application) → #2 Exploiting Server
- T1566.001 (Spearphishing Attachment) → #9 Social Engineering → #3 or #7 (depending on payload)
- T1078 (Valid Accounts) → #4 Identity Theft

Key distinction: Techniques describe attacker actions and behaviors (operational detail), while TLCTC clusters categorize the fundamental vulnerabilities being exploited (strategic framework). TLCTC proposes enhancing MITRE ATT&CK by adding cluster mappings and typical velocity attributes to each technique.

See also: TTP, Sub-Threat, MITRE ATT&CK, Operational Layer, Weakness

Temporal Notation (V2.0)

The V2.0 extension to standard attack path notation that explicitly annotates time intervals between threat cluster transitions (Δt) using the format $\rightarrow[\text{time}]$. Time units include seconds (s), minutes (m), hours (h), days, weeks, months. Examples:

- Basic: `#9→[24h]#4→[12m]#1`
- With domain boundaries: `#9→[days]#4→[mins]#1 || [dev][@Vendor→@Org] || →[weeks]#10.2→[0s]#7`
- With parallel execution: `#9→[30s]#7→[2m]#4→[15m](#1+#7)`

Enables precise velocity analysis, detection coverage score calculation, and realistic assessment of control effectiveness against time-sensitive attacks.

Transit Boundary Operator ($\Rightarrow$) (V2.1)

Notation: `||[context][@Source⇒@Carrier→@Target]||`. An extension to the Domain Boundary Operator that marks responsibility spheres which **carry or relay** the attack without being the source or the target. The $\Rightarrow$ symbol denotes transit (relay), while $\rightarrow$ denotes delivery to the final target. Chained transit uses right-to-left relay order: `||[context][@Source⇒@CarrierB⇒@CarrierA→@Target]||`. Transit is distinct from #10 Supply Chain Attack: transit marks a passive relay, while #10 marks a Trust Acceptance Event. Key rule (R-TRANSIT-3): vendor code running on the target device is NOT transit — it is the attack surface (classify by R-ROLE). Example: `#9 ||[human] [@Attacker⇒@SMSProvider→@Victim]||` — phishing SMS relayed through carrier.

Reference: §11.3.5 (Transit Boundary Operator)

Third-Party Trust Link (TTL)

Any reliance relationship where a third party can influence your domain. Examples: software components/libraries/dependencies, update/distribution channels, federation relationships (IdP/SP), managed control planes, SaaS admin consoles, signing/attestation/provenance chains, firmware/hardware supply chains, CI/CD pipeline integrations.

Reference: §4.2.2 (Global Definitions), §4.1 (#10 Definition)

Threat (in TLCTC)

An initiating force that exploits a generic vulnerability and can trigger the central event (Loss of Control), implemented as a set of tactics, techniques, and procedures (TTP) that attackers apply to provoke an event or incident. In TLCTC, threats are implemented as the 10 Top Level Cyber Threat Clusters, each defined by exactly one generic vulnerability. Threats are positioned on the cause side of the Bow-Tie model, distinct from vulnerabilities, events, and consequences (Axiom III). Threats are NOT outcomes, actors, or control failures.

Reference: §6.1 (Bow-Tie Structure), Axioms III–V (§2)

Threat Cluster

An organizational construct that groups a set of threats exploiting a common generic vulnerability related to IT systems and humans. The TLCTC framework defines 10 mutually exclusive threat clusters, each associated with a unique generic vulnerability (Axiom I) and distinct attack vector (Axiom II).

Threat Topology

A structural property of TLCTC describing whether a threat cluster (or a concrete attack step) operates primarily within the software domain's technical attack surfaces (**internal**) or enables crossing domain boundaries (**bridge**).

Reference: §5.0 (Topology in TLCTC), §5.1 (Definitions)

Tie-Breaker Rules

Precedence rules applied when a step appears to fit multiple clusters. Applied in order: (1) classify by initial generic vulnerability, (2) implementation flaw vs legitimate function misuse, (3) credential use always wins for the use step, (4) MitM starts at controlled position, (5) flooding is about capacity, (6) FEC execution must be explicit, (7) human/physical/third-party are not shortcuts, (8) document non-obvious decisions.

Reference: §4.2.6 (Tie-Breaker / Precedence Rules)

TEMPEST (*Industry Term*)

A codename for standards and techniques related to electromagnetic emanation security — both the interception of unintentional electromagnetic emissions from electronic equipment and the shielding against such interception. In TLCTC: TEMPEST attacks map to #8 Physical Attack (specifically the #8.2 signal vector, no contact) — exploiting physical emanation properties without direct contact with the target device. Related techniques include Van Eck Phreaking.

Reference: V1.9.1 Buzz-Word Refinement (#8)

See also: Physical Attack (#8), Van Eck Phreaking, RFID Skimming

Token Hijacking (*Industry Term*)

The theft or manipulation of authentication tokens (OAuth tokens, session tokens, API keys, bearer tokens) to gain unauthorized access. In TLCTC: the **acquisition** of the token maps to the enabling cluster per R-CRED (e.g., #3 via XSS, #5 via MitM interception, #7 via malware). The **use** of the stolen token to authenticate always maps to #4 Identity Theft.

Reference: V1.9.1 Buzz-Word Refinement (#4)

See also: Identity Theft (#4), Session Hijacking, Credential Application

Trojan (*Industry Term*)

Malware disguised as legitimate software to trick users into installing it. In TLCTC: the execution of the trojan maps to #7 Malware (foreign code running via designed execution capability). The deception that leads to installation typically involves #9 Social Engineering (tricking the user). Full sequence: #9 → #7. The trojan may then perform various malicious actions that map to additional clusters (e.g., credential stealing → #4, data exfiltration → [DRE: C]).

Reference: V1.9.1 Buzz-Word Refinement (#7)

See also: Malware (#7), Social Engineering (#9)

Typosquatting (*Industry Term*)

Registering domain names or package names that are slight misspellings of legitimate ones to deceive users or automated systems. In TLCTC: when used for malicious package publication (e.g., npm typosquatting), maps to #1 Abuse of Functions (abusing the publish functionality) → #10 Supply Chain Attack (trust boundary crossing) → #7 Malware (code execution at consumer). When used for phishing domains, see Domain Squatting.

Reference: V1.9.1 §Attack Path Notation (Example 3)

See also: Domain Squatting, Supply Chain Attack (#10), Abuse of Functions (#1)

TLCTC (Top Level Cyber Threat Clusters)

A pragmatic and structured framework for targeted threat identification that provides a universal approach to cybersecurity applicable across diverse IT systems and contexts. The framework consists of 10 distinct, non-overlapping threat clusters based on generic vulnerabilities, each with strategic and operational applications. The 10 Top Level Cyber Threat Clusters are:

1. Abuse of Functions
2. Exploiting Server
3. Exploiting Client
4. Identity Theft
5. Man in the Middle (MitM)

6. Flooding Attack
7. Malware
8. Physical Attack
9. Social Engineering
10. Supply Chain Attack

TLCTC Enumeration

A structured identifier system (TLCTC-XX.YY) where:

- TLCTC- prefix ensures proper attribution to the model
- XX represents the primary cluster number (01-10), zero-padded for consistent formatting
- .YY suffix designed for future refinement (.00 designates current high-level definitions)

This provides machine readability, consistent sorting, and extensibility for sub-categorization.

Related reading: [Dual-layer notation — TLCTC-XX.YY enumeration](#)

Trust Acceptance Event (TAE)

The moment your domain **honors** the Third-Party Trust Link and treats a Trust Artifact/Decision as authoritative. Actions at TAE include: validate, accept, install, apply, execute, attach privileges. #10 Supply Chain Attack is placed at the TAE.

Reference: §4.2.2 (Global Definitions), R-SUPPLY (§4.2.5), §4.1 (#10 Definition)

Trust Artifact / Trust Decision (TAD)

What crosses the boundary and is accepted as authoritative in a third-party trust relationship. Examples: SAML/OIDC assertions, federated tokens, signed packages/updates/container images, CI build artifacts/release binaries, policy/configuration pushes, admin actions from managed platforms, firmware images.

Reference: §4.2.2 (Global Definitions), §4.1 (#10 Definition)

TTP (Tactics, Techniques, and Procedures)

A detailed description of attacker behavior. In the TLCTC framework, specific TTPs (like those in MITRE ATT&CK) are considered instances or sub-threats that are categorized under the broader, cause-oriented Top Level Cyber Threat Clusters.

Two-Tiered Approach

The TLCTC structure distinguishing between:

- **Strategic Management Layer:** High-level risk management, policy-making, and governance using the 10 Top Level Cyber Threat Clusters
- **Operational Layer:** Detailed implementation of controls, specific vulnerability management, and threat intelligence using sub-threats and TTPs

U

UDP Flood (*Industry Term*)

A network-layer denial of service attack that overwhelms a target with UDP packets, consuming bandwidth and processing resources. In TLCTC: maps to #6 Flooding Attack — volume exceeding finite network/processing capacity.

Reference: V1.9.1 Buzz-Word Refinement (#6)

See also: Flooding Attack (#6), DDoS, SYN Flood

Umbrella Controls

Security measures that provide protection for groups of IT systems within their scope, such as firewalls, proxies, network zones, or external network filters. These contrast with Local Controls that protect specific systems directly. Important consideration: Umbrella controls provide protection only for specific 'Groups of IT-Systems' within their scope and cannot effectively protect all exposed systems (e.g., a firewall protects 'inner IT-Systems' but not directly exposed ones).

Unknown Δt

A Δt value where no supported time statement can be made. Notation: $\Delta t=?$.

Reference: §12.0.3, §12.2.3

Unresolved-Step Operators (?, ...) (V2.1)

Notation operators for partially-resolved attack paths where forensic evidence confirms that a step (or gap of steps) exists but the cluster cannot yet be determined. ? represents exactly one unresolved step; ... (or ASCII ...) represents a gap of one or more steps. Governed by the seven canonical R-UNRES rules (2, 3, 5, 6, 7, 8, 9 — the numbering is intentionally non-contiguous; draft rules 1 and 4 were consolidated into adjacent rules during v2.1 finalization). Key constraints: unresolved steps MUST NOT carry DRE tags (R-UNRES-5); if any cluster can be defended even weakly, the step MUST be classified as #X [conf=low] rather than left unresolved (R-UNRES-9); every unresolved step MUST be accompanied by a prose annotation (R-UNRES-8).

Reference: §11.5.4 (Unresolved-Step Operators)

USB Baiting (*Industry Term*)

A physical attack where an attacker leaves malicious USB devices in locations where targets are likely to find and connect them (parking lots, lobbies, conference rooms). In TLCTC: the physical placement of the USB device maps to #8 Physical Attack. If the victim is induced by curiosity to plug it in, social engineering is also involved (#9). The execution of malicious code from the device maps to #7 Malware. Typical sequence: #8 → #7 or #9 → #8 → #7.

Reference: V1.9.1 Buzz-Word Refinement (#8)

See also: Physical Attack (#8), Malware (#7), Evil Maid Attack

V

Velocity Annotation

Notation: $\rightarrow[\Delta t=\text{value}]$ or $\rightarrow[\Delta t=Xh]$, $\rightarrow[\Delta t=Xm]$, $\rightarrow[\Delta t=Xs]$. Indicates the observed or estimated time interval between one Attack Step and the next. Velocity annotations are OPTIONAL but RECOMMENDED for operational analysis and threat intelligence sharing.

Reference: §4.2.2 (Global Definitions), §11.5.2, §12.2 (Δt Notation)

Velocity Class

Categorical labels for Δt ranges that describe the defender's feasible response mode and determine appropriate control strategies. Four classes are defined:

- **VC-1: Strategic / Latent/Slow** (days $\rightarrow$ months): Log retention, threat hunting, strategic monitoring
- **VC-2: Tactical / Medium** (hours): SIEM alerting, analyst triage, guided response
- **VC-3: Operational / Fast** (minutes): Automation (SOAR/EDR), rapid containment, prebuilt playbooks
- **VC-4: Real-Time** (seconds $\rightarrow$ milliseconds): Architecture & circuit breakers, rate-limits, hardening, automatic isolation

Reference: §12.4 (Operational Velocity Classes)

Van Eck Phreaking (*Industry Term*)

A technique for eavesdropping on the contents of a CRT or LCD display by detecting and decoding the electromagnetic emissions produced by the display. In TLCTC: maps to #8 Physical Attack (specifically the #8.2 signal vector, no contact) — exploiting electromagnetic emanation properties without direct physical contact.

Reference: V1.9.1 Buzz-Word Refinement (#8)

See also: Physical Attack (#8), TEMPEST

Vertical Stack Application

The implementation of TLCTC across the layered architecture of IT systems (from application level to hardware), analyzing client/server roles at each protection ring boundary (e.g., Ring 3 to Ring 0) and directional vulnerabilities.

Vishing (*Industry Term*)

Social engineering attacks delivered via voice calls (phone). In TLCTC: maps to #9 Social Engineering — the delivery channel (voice) does not change the cluster classification. Subsequent steps map to their respective clusters.

Reference: V1.9.1 Buzz-Word Refinement (#9)

See also: Social Engineering (#9), Phishing, Smishing

Vulnerability

An exploitable condition in a system that constitutes the **attack surface towards a threat**. A vulnerability is what the threat "sees" and targets — it is the exposed surface that enables the attack vector and, through it, the attack path. Without the vulnerability, the threat has no point of entry; with it, the threat has a viable route from cause to compromise.

In the TLCTC Bow-Tie model this relationship is structural: each of the 10 threat clusters is defined by exactly one **generic vulnerability** (Axiom I). The generic vulnerability is the attack surface that the threat cluster exploits. It is what makes the attack vector possible and what the attack path traverses. Controls exist to reduce or eliminate this exposed surface — preventive controls on the left side of the Bow-Tie shrink the attack surface; the vulnerability that remains is the residual exposure the threat can still reach.

Conceptual hierarchy: Weakness (CWE) → Specific Vulnerability (CVE) → Generic Vulnerability (TLCTC) → Threat Cluster (#1–#10).

- A **weakness** (CWE) is the underlying flaw, bug, or error that creates the condition.
- A **specific vulnerability** (CVE) is a concrete, exploitable instance of that condition in a particular product or version.
- A **generic vulnerability** (TLCTC) is the universal, technology-independent category of attack surface that all specific vulnerabilities of the same nature map to.
- A **threat cluster** is the set of threats that target that generic attack surface.

Example: A coding error that fails to validate input (weakness / CWE-89) creates a SQL injection vulnerability (specific vulnerability / CVE-xxxx-yyyy) in a web application. That vulnerability is the attack surface towards #2 Exploiting Server — the exposed surface through which the attacker's exploit code enters. The attack vector is defined by this initial generic vulnerability (Axiom VII), and the attack path proceeds from there (e.g., #2 → #7 → #4).

Critical distinction: A vulnerability is an exploitable condition that exists in a system, while a weakness is the underlying flaw, bug, or error that enables that vulnerability to exist. TLCTC focuses on categorizing the generic vulnerabilities — the fundamental attack surfaces — that all specific vulnerabilities map to.

See also: Generic Vulnerability, Weakness, CVE, Threat Cluster, Attack Vector, Attack Path, Bow-Tie Model

W

WAF (Web Application Firewall) (*Industry Term*)

A security control that monitors, filters, and blocks HTTP traffic to and from a web application. In TLCTC: WAF is an **Umbrella Control** primarily targeting #2 Exploiting Server — it provides a protective layer against common server-side exploitation techniques (SQL injection, XSS, etc.) for web applications within its scope. WAFs cannot protect against all threat clusters and should be part of a defense-in-depth strategy.

See also: Umbrella Controls, Exploiting Server (#2), Defense-in-Depth

Watering Hole Attack (*Industry Term*)

An attack where an adversary compromises a website frequently visited by the target group, then uses the compromised site to deliver exploits or malware to visitors. In TLCTC: the attack decomposes into multiple steps depending on the method:

- Compromising the website: maps to its own cluster (e.g., #2 Exploiting Server, #4 Identity Theft)
- Luring users to visit: #9 Social Engineering (if actively lured) or implicit via the site's normal traffic
- Exploiting visitor's browser: #3 Exploiting Client
- Delivering malware: #7 Malware

Typical sequence: #3 → #7 (for the victim's perspective) or the full chain from attacker's perspective.

Reference: V1.9.1 Buzz-Word Refinement (#3)

See also: Exploiting Client (#3), Drive-By Download, Malvertising

Weakness

A flaw, bug, or error in software, hardware, or processes that enables vulnerabilities to exist. In the Common Weakness Enumeration (CWE) framework, weaknesses are categorized as the root causes of software security problems (e.g., CWE-89 for SQL Injection weakness, CWE-119 for buffer overflow weakness).

Critical distinction in TLCTC context: CWE categorizes weaknesses (the flaws themselves), not vulnerabilities (the exploitable conditions those flaws create). In the TLCTC framework, the conceptual hierarchy flows: **Weakness** → **Specific Vulnerability (CVE)** → **Generic Vulnerability** → **Threat Cluster**.

Example: A coding error that fails to validate input (weakness) creates a SQL injection vulnerability (specific vulnerability), which exploits the generic vulnerability of "server-side code flaws" (#2 Exploiting Server). TLCTC's 10 generic vulnerabilities represent the universal categories that all specific vulnerabilities ultimately map to, regardless of their underlying weaknesses.

Relationship to TLCTC: While CWE provides granular weakness taxonomy at the code level for developers, TLCTC operates at the strategic level by grouping all resulting vulnerabilities into 10 generic vulnerability categories that define the threat clusters. Both frameworks are complementary.

See also: Vulnerability, Generic Vulnerability, CVE, CWE, Threat Cluster, Coder, Programmer

Whaling (*Industry Term*)

A targeted phishing attack aimed specifically at senior executives or high-value targets within an organization. In TLCTC: maps to #9 Social Engineering — the seniority of the target is an operational detail that does not change the cluster classification. Often used in BEC (Business Email Compromise) / CEO Fraud scenarios. Subsequent steps map to their respective clusters.

Reference: V1.9.1 Buzz-Word Refinement (#9)

See also: Social Engineering (#9), Spear Phishing, Phishing

Worm (*Industry Term*)

Self-replicating malware that spreads across networks without requiring user interaction, typically by exploiting vulnerabilities in network-accessible services. In TLCTC: the exploitation of the vulnerability maps to #2 Exploiting Server (or #3 Exploiting Client), and the execution of the worm payload maps to #7 Malware. Wormable exploits are characteristic of the **Realtime Velocity Class** (seconds/milliseconds) — e.g., EternalBlue — where architecture and hardening are the only effective controls because detection/response cannot keep pace.

Reference: V1.9.1 Buzz-Word Refinement (#7), Realtime Velocity Class

See also: Malware (#7), Exploiting Server (#2), Realtime Velocity Class

X

XXE (XML External Entity) Injection (*Industry Term*)

An implementation flaw where an application processes XML input containing references to external entities, potentially leading to data disclosure, SSRF, or denial of service. In TLCTC: maps to #2 Exploiting Server (or #3 Exploiting Client per R-ROLE) — a coding flaw in XML parsing that creates an unintended data→code transition.

Reference: V1.9.1 Buzz-Word Refinement (#2)

See also: Exploiting Server (#2), SSRF, Implementation Flaw

Quick Reference Tables

Cluster Quick Reference

#	Name	Generic Vulnerability	Topology
#1	Abuse of Functions	Functional scope/trust (designed capabilities abused)	Internal
#2	Exploiting Server	Server-side code implementation flaws	Internal
#3	Exploiting Client	Client-side code implementation flaws	Internal
#4	Identity Theft	Insufficient identity-artifact binding at the point of authentication	Internal
#5	Man in the Middle	Lack of end-to-end communication protection	Internal
#6	Flooding Attack	Finite capacity limitations	Internal
#7	Malware	Designed execution capability for untrusted content	Internal
#8	Physical Attack	Physical accessibility/interference	Bridge
#9	Social Engineering	Human psychological factors	Bridge
#10	Supply Chain Attack	Third-party trust dependencies	Bridge

Reference: §4.1 (Cluster Definitions), §5.2 (Topology Classification)

Related reading: [Mandiant M-Trends 2025 — TLCTC](#), [ENISA Threat Landscape 2025 — TLCTC](#), [Same Attack, Four Stories — vendor report comparison](#), [TLCTC × Threat Modeling Manifesto](#), [End of Semantic Diffusion — Kuhn & TLCTC](#), [TLCTC classification decision tree V2.0/V2.1](#), [Logical foundations of TLCTC](#), [Why exactly ten? — TLCTC architecture](#), [10×10×10 — fun fact about TLCTC](#), [Agentic AI under the microscope — TLCTC](#), [Quantum & AI — new magic, same 10 threats](#), [AI conversation deep dive on TLCTC](#), [Grok AI 4.1 — independent TLCTC validation](#), [G7 SBOM-for-AI — control fixation critique](#), [Capability-based planning via 10×\(6×2\) matrix](#), [TLCTC — the missing link between strategy and ops](#), [Tactics evolve, 10 threats are constant](#), [Tactics evolve, 10 threats are constant — TLCTC](#), [TLCTC for everyone — the Blind Spot method](#), [Emerging technologies as threat enablers](#), [TLCTC v2.5 monster prompt — CISO & Risk](#), [TLCTC v2.5 monster prompt — CTI & Forensic](#), [TLCTC v2.5 monster prompt — Regulators & Standards](#), [TLCTC v2.1 AI analysis prompt — teach any LLM](#), [TLCTC v2.5 monster prompts — index](#), [The Risk Appetite Conversation Your Board Isn't Having](#), [Understanding cyber threats — a common language](#), [TLCTC — the missing link \(brief\)](#), [Strategic risk management implementation guide v2.1](#)

Axiom Quick Reference

#	Name	Group	Core Statement
I	No System-Type Differentiation	Scope	Generic IT assets; sector labels don't create threat classes
II	Client–Server Model	Scope	Universal interaction abstraction
III	Causes, Not Outcomes	Separation	Threats ≠ data risk events
IV	Not Threat Actors	Separation	Threats ≠ actor identity
V	Not Control Failure	Separation	Control-risk ≠ threat category
VI	Single-Cluster Rule	Classification	One step = one vulnerability = one cluster
VII	Initial-Vulnerability Rule	Classification	Vector defined by initial generic vulnerability
VIII	Strategic–Operational Layering	Classification	Clusters → sub-threats
IX	Sequence + Velocity	Sequences	Clusters chain; Δt measures velocity
X	Credential Duality	Sequences	Acquisition vs application

Reference: §2 (Axioms and Assumptions)

R-* Rules Quick Reference

The v2.5 normative registry contains exactly **19 rules**: nine core rules and ten v2.1 extension rules. This table mirrors the canonical dictionary (`tlctc-framework.v2.5.json`); summaries are condensed, the dictionary statement governs.

Core rules (9):

Rule	Distinguishes	Key Decision
R-EXEC	FEC Execution	If FEC executes → #7 MUST be recorded (plus enabling cluster)
R-ROLE	#2 vs #3	Server-role (accepts inbound) → #2 ; Client-role (consumes external) → #3 ; roles set by call direction at any interface, network not required
R-FLOOD	Capacity vs Defect	Volume exhaustion → #6 ; Implementation defect → #2/#3 per R-ROLE
R-SUPPLY	TAE Placement	#10 at Trust Acceptance Event where third-party trust is honored
R-MITM	Position vs Action	Gaining position → enabling cluster; Interception/modification/relay → #5
R-CHANNEL (v2.5)	Control vs Code Flaw	Defective logic constitutive of channel control → #5 ; Incidental defect → #2/#3
R-SUBSTRATE (v2.5)	Property vs Logic	Physical property exploited → #8 ; Physical layer as readout only → #2/#3
R-CRED	Acquisition vs Use	Acquisition → enabling cluster; Use → always #4 (unless the identity is the presenter's own — self-issued enrolment is #1); separate steps
R-SCOPE (v2.5)	Attack vs Abuse of Rights	Actor? Intent? Entitlement covering this action? → in-grant intended action = Abuse of Rights (OpRisk, no cluster, no SRE); unentitled → Attack row, clusters apply

v2.1 extension rules (10):

Rule	Distinguishes	Key Decision
R-TRANSIT-3 (V2.1)	Transit vs Attack Surface	Vendor code on target device → classify by R-ROLE, not transit
R-INTRA-7 (V2.1)	Annotation vs Classification	Intra-system crossings never change cluster classification
R-INTRA-9 (V2.1)	Memory Deferral	memory boundary type deferred → MUST NOT use
R-UNRES-2 (V2.1)	Annotation vs Cluster	? / ... are epistemic annotations, not clusters (no #11 / #12)
R-UNRES-3 (V2.1)	Statistics	? / ... excluded from statistics
R-UNRES-5 (V2.1)	No DRE on Unresolved	DRE tags MUST NOT be appended to ? or ...
R-UNRES-6 (V2.1)	Boundary Compatibility	Boundary operators MAY appear adjacent to ? / ...
R-UNRES-7 (V2.1)	Resolution Obligation	Every ? / ... is an open analytical task → resolve when evidence arrives
R-UNRES-8 (V2.1)	Prose Required	Paths containing ? / ... MUST have prose annotation explaining gap
R-UNRES-9 (V2.1)	Classification Threshold	If any cluster can be defended → use #X [conf=low] , not ?

Deprecated / withdrawn rule IDs (do not cite as normative):

ID	Status	Where the substance lives now
R-ABUSE	Deprecated v2.0 alias	#1 cluster definition and boundary tests (core §4)
R-HUMAN	Deprecated v2.0 alias	#9 cluster definition and boundary tests (core §4)
R-PHYSICAL	Deprecated v2.0 alias	#8 boundary tests (core §4) + R-SUBSTRATE (admission)
R-TRANSIT-1/2/4/5/6/7/8	Withdrawn v2.1 draft series	Non-normative transit notation practice (see R-TRANSIT entry); cluster independence is SG-2
R-INTRA-1...6, -8	Withdrawn v2.1 draft series	Non-normative intra-system notation practice (see R-INTRA entry)
R-UNRES-1, -4	Consolidated during v2.1 finalization	R-UNRES-1 → R-UNRES-2; draft R-UNRES-4's threshold → canonical R-UNRES-9

Caution: the draft series numbered some propositions differently than the canon (draft R-INTRA-7 = "distinct vulnerabilities", canonical R-INTRA-7 = "no cluster change"; draft R-UNRES-4 = the conf=low threshold, canonical R-UNRES-9). Always cite the canonical registry above.

Semantic Guardrails Quick Reference (V2.1)

ID	Rule	Key Constraint
SG-1	Cause First	Classify by generic vulnerability exploited, not topology or effect
SG-2	Topology ≠ Classification	Transit/intra-system annotations MUST NOT define or imply a cluster
SG-3	Annotations Subordinate	Annotations MUST NOT appear as independent path elements or replace clusters
SG-4	Effects ≠ Threats	Sandbox escape, privilege escalation, etc. are not TLCTC clusters
SG-5	Actors ≠ Threats	Transit parties, vendors, carriers are spheres, not threat categories
SG-6	Distinct Exploit Rule	Separately evidenced exploit → new cluster step required
SG-7	Backward Recoverability	Stripping V2.1 annotations MUST leave a valid cluster sequence

Reference: §4.2.4 (Semantic Guardrails), §4.2.5 (Global Mapping Rules)